

# M2177.0043 Introduction to Deep Learning

## Lecture 5: Optimization<sup>1</sup>

Hyun Oh Song<sup>1</sup>

<sup>1</sup>Dept. of Computer Science and Engineering, Seoul National University

March 31, 2020

---

<sup>1</sup>Many slides and figures adapted Justin Johnson

## Last time

- ▶ Constrained optimization
- ▶ Online method

# Outline

Subgradient

Online method

# Subgradient

Not all functions are differentiable. The absolute value  $f(x) = |x|$ , for example, is convex but not differentiable at 0.

## Definition 1 (Subgradients)

Vector  $g$  is a **subgradient** of  $f$  at  $x \in \text{dom} f$  if


$$f(y) \geq f(x) + g^\top (y - x) \quad \forall y \in \text{dom} f$$

The set of all subgradients at  $x$  is called the **subdifferential** at  $x$  and is denoted as  $\partial f(x)$ .

The existence of subgradients is often sufficient for optimization.

## Examples<sup>2</sup>

Consider  $f : \mathbb{R} \rightarrow \mathbb{R}, f(x) = |x|$




- ▶ For  $x \neq 0, |y| \geq |x| + g(y - x), \forall y$ . Or  $|y| - gy \geq |x| - gx$ . Satisfies if  $|x| - gx = 0$ . This means  $g = \text{sign}(x)$ .
- ▶ For  $x = 0$ , inequality simplifies to  $|y| \geq gy$ . This means  $g \in [-1, 1]$ .

---


<sup>2</sup>taken from Tibshirani 10-725

Consider  $f : \mathbb{R}^n \rightarrow \mathbb{R}, f(x) = \|x\|_2$




- ▶ For  $x \neq 0$ , differentiable  $g = \frac{x}{\|x\|}$
- ▶ For  $x = 0$ , inequality simplifies to  $\|y\|_2 \geq g^\top y$ . So  $\|y\|_2 \geq \|g\|_2 \|y\|_2 \cos \theta$ .  $g \in \{z \mid \|z\|_2 \leq 1\}$

Consider  $f : \mathbb{R}^n \rightarrow \mathbb{R}, f(x) = \|x\|_1$



- ▶ For  $x \neq 0$ , differentiable  $g_i = \text{sign}(x_i)$
- ▶ For  $x = 0$ , inequality simplifies to  $\|y\|_1 \geq g^\top y$ . So  $g_i \in [-1, 1]$ .

Let  $f_1, f_2 : \mathbb{R}^n \rightarrow \mathbb{R}$ , be convex differentiable and consider  
 $f(x) = \max\{f_1(x), f_2(x)\}$



- ▶ For  $f_1(x) > f_2(x)$ ,  $g = \nabla f_1(x)$
- ▶ For  $f_2(x) > f_1(x)$ ,  $g = \nabla f_2(x)$
- ▶ For  $f_1(x) = f_2(x)$ , subgradient  $g$  is any point on the line segment between  $\nabla f_1(x)$  and  $\nabla f_2(x)$ . Concretely,  
 $g \in \{\theta_1 \nabla f_1(x) + \theta_2 \nabla f_2(x) \mid \theta_1 + \theta_2 = 1, \theta_1 \geq 0, \theta_2 \geq 0\}$



# Outline

Subgradient

Online method

## Stochastic gradient descent

- Consider minimizing an objective function that has the form of a sum of functions:

$$f(x) = \frac{1}{n} \sum_{i=1}^n f_i(x)$$

- Each summand function  $f_i(x)$  is typically associated with the  $i$ -th observation in a dataset. The standard (or “batch”) gradient descent method would perform the following iterations:

$$x := x - \eta \nabla f(x) = x - \eta \sum_{i=1}^n \frac{1}{n} \nabla f_i(x)$$

- Computing the gradient can be very expensive if  $n$  is large. Stochastic gradient descent method *samples* the summand functions at every step for scalability.

$$x := x - \eta \nabla f_i(x)$$

## Sanity check<sup>3</sup>

- ▶ Let us check that on a simple problem that the stochastic gradient descent yields the optimum. Let  $p_1, \dots, p_m \in \mathbb{R}^n$ , and define  $f: \mathbb{R}^n \rightarrow \mathbb{R}_+$ :

$$\forall x \in \mathbb{R}^n, f(x) = \frac{1}{2m} \sum_{i=1}^m \|x - p_i\|_2^2$$

- ▶ Note that here  $f_i(x) = \frac{1}{2} \|x - p_i\|_2^2$  and  $\nabla f_i(x) = x - p_i$ . Moreover,

$$x^* = \operatorname{argmin}_{x \in \mathbb{R}^d} f(x) = \frac{1}{m} \sum_{i=1}^m p_i$$

---

<sup>3</sup>Taken from Moritz Hardt's lecture

- Now, run SGM with  $\eta_t = \frac{1}{t}$  in cyclic order i.e.  $i_t = t$  and  $x_0 = 0$ :

$$x_0 = 0$$

$$x_1 = 0 - \frac{1}{1}(0 - p_1) = p_1$$

$$x_2 = p_1 - \frac{1}{2}(p_1 - p_2) = \frac{p_1 + p_2}{2}$$

$$\vdots$$

$$x_m = \frac{1}{m} \sum_{i=1}^m p_i = x^*$$

## Minibatch SGD

- ▶ A compromise between computing the true gradient and the stochastic gradient of a single example
- ▶ Reduces the variance of the gradient estimate
- ▶ In practice, vector-process the gradient computation so the minibatch size fits in the memory

## Minibatch SGD

---

*Minibatch stochastic gradient descent*

**given** a starting point  $x \in \text{dom} f$ .

**repeat**

1. *Shuffle the data*
2. **For every m sequence of data**  $i = 1, \dots, \lceil n/m \rceil$ 
  - a. *Compute the minibatch gradient.*  $\Delta x = \frac{1}{m} \sum_{j=1}^m \nabla f_j(x)$
  - b. *Update the stepsize.*  $t := \text{Update}(t)$ .
  - c. *Update.*  $x := x + t\Delta x$ .
3. *Decay the stepsize*

**until** stopping criterion is satisfied.

---

## Distributed gradient descent for full batch gradient descent

- ▶ **Map:** compute gradient on subblock and emit
- Reduce:** aggregate parts of the gradients



---

**given** a starting point  $x \in \text{dom} f$ .

**repeat**


1. *Compute the full gradient*  $\Delta x_{\text{ds}} := \sum_{i=1}^n \frac{1}{n} \nabla f_i(x)$ .
2. *Line search.* Choose a step size  $t > 0$ .
3. *Update.*  $x := x + t \Delta x_{\text{ds}}$ .

**until** stopping criterion is satisfied.

---

## Problems in non-convex optimization

- ▶ What if the loss function has a local minima or saddle point?
- ▶ Zero gradient, gradient descent gets stuck





## Momentum

- ▶ Think about roll a ball down a hill
- ▶ Build up “velocity” as a running mean of gradients

$$v_{t+1} = \rho v_t + \nabla f(x_t)$$

$$x_{t+1} = x_t - \alpha v_{t+1}$$

Typically set  $\rho = 0.9$  or  $0.99$ .



Figure: (Left) local minima, (Right) Saddle points

## AdaGrad

- ▶ Adapt the learning rate to each parameter dimensions. Larger updates for infrequent and smaller updates for frequent parameters.
- ▶ Keep a diagonal matrix  $G_t \in \mathbb{R}^{d \times d}$  where each diagonal element is the sum of squares of the gradients up to time step  $t$ .

$$x_{t+1} = x_t - \frac{\eta}{\sqrt{G_t + \epsilon}} \odot g_t$$

- ▶  $\epsilon$  is a small constant ( $\approx 1e-8$ ) to avoid divide by zero.

```
grad_squared = 0
while True:
    dx = compute_gradient(x)
    grad_squared += dx * dx
    x -= learning_rate * dx / (np.sqrt(grad_squared) + 1e-7)
```

- ▶ What happens to the denominator if you keep adding positive values?

## RMSProp

- ▶ As opposed to accumulating all previous squared gradients, restrict the attention to recent gradients of some time window. Divide the learning rate by an exponentially decaying average of squared gradients
- ▶ Default  $\gamma = 0.9, \eta = 0.001$  (by Geoff Hinton)

AdaGrad

```
grad_squared = 0
while True:
    dx = compute_gradient(x)
    grad_squared += dx * dx
    x -= learning_rate * dx / (np.sqrt(grad_squared) + 1e-7)
```



RMSProp

```
grad_squared = 0
while True:
    dx = compute_gradient(x)
    grad_squared = decay_rate * grad_squared + (1 - decay_rate) * dx * dx
    x -= learning_rate * dx / (np.sqrt(grad_squared) + 1e-7)
```

# Adam

- Combine momentum and AdaGrad/RMSProp

```
first_moment = 0
second_moment = 0
while True:
    dx = compute_gradient(x)
    first_moment = beta1 * first_moment + (1 - beta1) * dx
    second_moment = beta2 * second_moment + (1 - beta2) * dx * dx
    x -= learning_rate * first_moment / (np.sqrt(second_moment) + 1e-7))
```


Momentum

AdaGrad / RMSProp

Sort of like RMSProp with momentum


- What happens at first time step?

## Learning rate (step size) as a hyperparameter



Q: Which one of these learning rates is best to use?

## Learning rate (step size) as a hyperparameter



**=> Learning rate decay over time!**

**step decay:**

e.g. decay learning rate by half every few epochs.

**exponential decay:**

$$\alpha = \alpha_0 e^{-kt}$$

**1/t decay:**

$$\alpha = \alpha_0 / (1 + kt)$$

## Second-order optimization

- ▶ Quasi-Newton methods (BFGS most popular) Instead of inverting the Hessian ( $O(n^3)$ ), approximate inverse Hessian with rank 1 updates over time ( $O(n^2)$ ).
- ▶ L-BFGS (Limited memory BFGS) Does not form/store the full Hessian (would need  $O(n^2)$  space). Instead, only retain recent  $m$  gradients (around 10 to 20). Hence the name *Limited memory*.
- ▶ Usually works very well in full batch. Does not transfer well to mini-batch setting. Gives bad results. Adapting second order optimization to large scale stochastic setting is an active area of research.

## In practice

- ▶ **Adam** is a good default choice in most cases
- ▶ If you can afford to do full batch updates then try out **L-BFGS**.