

# M2177.0043 Introduction to Deep Learning

## Lecture 14: Deep Generative Models<sup>1</sup>

Hyun Oh Song<sup>1</sup>

<sup>1</sup>Dept. of Computer Science and Engineering, Seoul National University

May 7, 2020

---

<sup>1</sup>Many slides and figures adapted Justin Johnson

## Last time

- ▶ VAE
- ▶ Reparameterization trick

# Outline

Autoregressive models: PixelRNN and PixelCNN

Generative adversarial networks

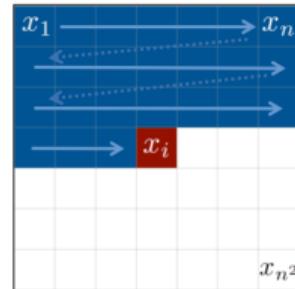
Summary of deep generative models

# PixelRNN<sup>2</sup>

## Autoregressive Image Modeling

- Autoregressive models train a network that models the conditional distribution of every individual pixel given previous pixels (raster scan order dependencies).

$$p(\mathbf{x}) = \prod_{i=1}^{n^2} p(x_i | x_1, \dots, x_{i-1}).$$



⇒ **sequentially predict pixels** rather than predicting the whole image at once (like as GAN, VAE)

- For color image, 3 channels are generated successive conditioning, **blue** given **red** and **green**, **green** given **red**, and **red** given only the pixels above and to the left of all channels.



<sup>2</sup>from Yohei Sugawara

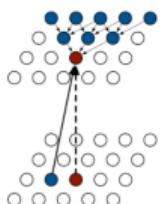
# PixelRNN<sup>3</sup>

## Previous work: Pixel Recurrent Neural Networks.

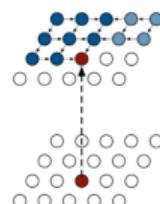
- “Pixel Recurrent Neural Networks” got best paper award at ICML2016.
- They proposed two types of models, **PixelRNN** and **PixelCNN**  
(two types of LSTM layers are proposed for PixelRNN.)

### PixelRNN

#### Row LSTM

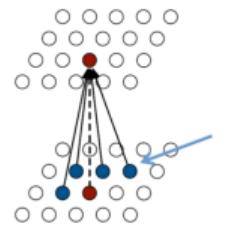


#### Diagonal BiLSTM



### PixelCNN

#### masked convolution



$$\begin{matrix} W_{11} & W_{12} & W_{13} \\ W_{21} & W_{22} & W_{23} \\ W_{31} & W_{32} & W_{33} \end{matrix}$$

$$\begin{matrix} 1 & 1 & 1 \\ 1 & 0 & 0 \\ 0 & 0 & 0 \end{matrix}$$

- LSTM based models are natural choice for dealing with the autoregressive dependencies.

- CNN based model uses masked convolution, to ensure the model is causal.

|       | PixelRNN   | PixelCNN  |
|-------|--|---|
| Pros. | <ul style="list-style-type: none"><li>effectively handles long-range dependencies<br/>⇒ <b>good performance</b></li></ul>            | Convolutions are easier to parallelize ⇒ <b>much faster</b> to train  |
| Cons. | <ul style="list-style-type: none"><li>Each state needs to be computed sequentially.<br/>⇒ <b>computationally expensive</b></li></ul> | <ul style="list-style-type: none"><li>Bounded receptive field ⇒ <b>inferior performance</b></li><li><b>Blind spot problem</b> (due to the masked convolution) needs to be eliminated.</li></ul> |

<sup>3</sup>from Yohei Sugawara

# PixelRNN

- ▶ Refer to the original paper  
<https://arxiv.org/abs/1601.06759> for the implementation details and parallelizations

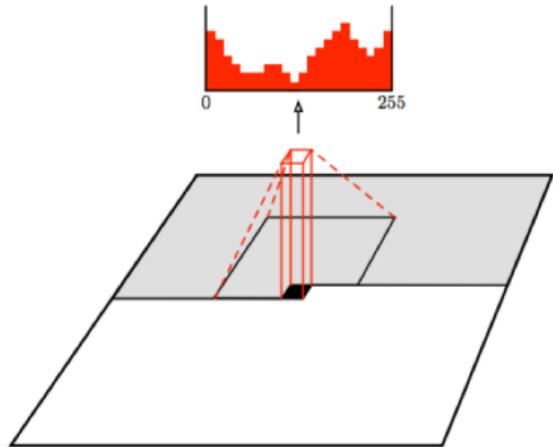
# PixelCNN

- ▶ Dependency on previous pixels now models using a CNN with **masked convolutions** over context region
- ▶ Training: maximize likelihood of training images

$$p(x) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

- ▶ Training faster than PixelRNN by parallelizing convolutions over each context regions
- ▶ Generation still slow because it need to happen sequentially

Softmax loss at each pixel



# Outline

Autoregressive models: PixelRNN and PixelCNN

Generative adversarial networks

Summary of deep generative models

## Generative adversarial networks

- ▶ VAEs define intractable density function with latent variable  $z$ .

$$p_{\theta}(x) = \int p_{\theta}(z)p_{\theta}(x | z)dz$$

However, we cannot optimize this directly, so we derive and optimize the lower bound on likelihood (ELBO) instead

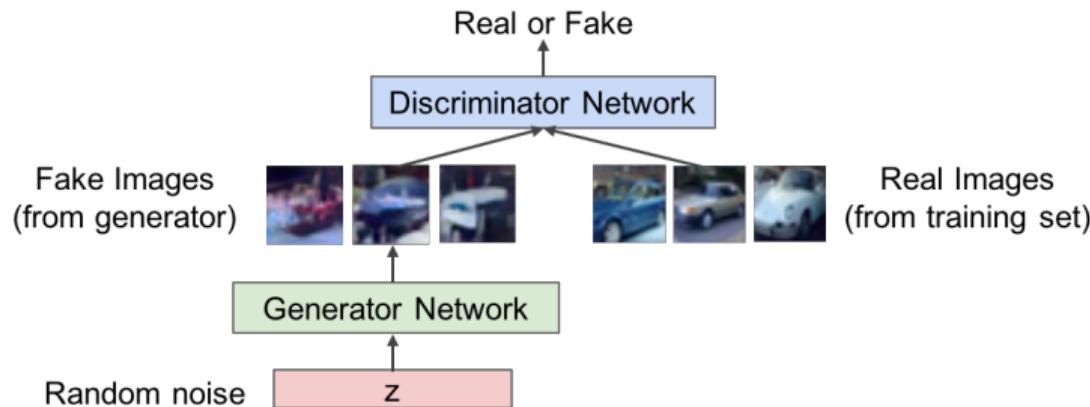
- ▶ GANs: don't work with any explicit density function. Instead, take game-theoretic approach: learn to generate from training distribution through 2-player game.

## Generative adversarial networks

- ▶ Problem: want to sample from complex, high-dimensional training distribution. No direct way to do this!
- ▶ Solution: Sample from a simple distribution. e.g. random noise. Learn transformation (generator network) to training distribution.

## Training GANs: Two-player game

- ▶ Generator network: try to fool the discriminator by generating real-looking images
- ▶ Discriminator network: try to distinguish between real and fake images



Fake and real images copyright Emily Denton et al. 2015. Reproduced with permission.

## Training GANs: Two-player game

- ▶ Generator network: try to fool the discriminator by generating real-looking images
- ▶ Discriminator network: try to distinguish between real and fake images

Train jointly in **minimax game**

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

## Training GANs: Two-player game

- ▶ Generator network: try to fool the discriminator by generating real-looking images
- ▶ Discriminator network: try to distinguish between real and fake images

Train jointly in **minimax game**

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log \underbrace{D_{\theta_d}(x)}_{\text{Discriminator output for real data } x} + \mathbb{E}_{z \sim p(z)} \log \underbrace{(1 - D_{\theta_d}(G_{\theta_g}(z)))}_{\text{Discriminator output for generated fake data } G(z)} \right]$$

## Training GANs: Two-player game

- ▶ Generator network: try to fool the discriminator by generating real-looking images
- ▶ Discriminator network: try to distinguish between real and fake images

Train jointly in **minimax game**

Discriminator outputs likelihood in (0,1) of real image

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log \underbrace{D_{\theta_d}(x)}_{\text{Discriminator output for real data } x} + \mathbb{E}_{z \sim p(z)} \log(1 - \underbrace{D_{\theta_d}(G_{\theta_g}(z))}_{\text{Discriminator output for generated fake data } G(z)}) \right]$$

- Discriminator ( $\theta_d$ ) wants to **maximize objective** such that  $D(x)$  is close to 1 (real) and  $D(G(z))$  is close to 0 (fake)
- Generator ( $\theta_g$ ) wants to **minimize objective** such that  $D(G(z))$  is close to 1 (discriminator is fooled into thinking generated  $G(z)$  is real)

## Training GANs: Two-player game

Train jointly in **minimax game**

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

## Training GANs: Two-player game

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. **Gradient ascent** on discriminator

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. **Gradient descent** on generator

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

# Training GANs: Two-player game

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. Gradient ascent on discriminator

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

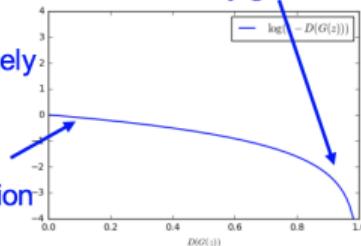
Gradient signal dominated by region where sample is already good

2. Gradient descent on generator

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

In practice, optimizing this generator objective does not work well!

When sample is likely fake, want to learn from it to improve generator. But gradient in this region is relatively flat!



x-axis:  $D(G(z))$ , y-axis (blue):  $\log(1 - D(G(z)))$

# Training GANs: Two-player game

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. Gradient ascent on discriminator

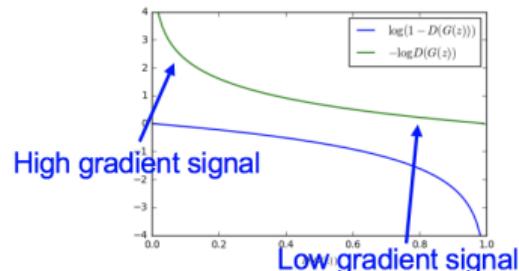
$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. Instead: Gradient ascent on generator, different objective

$$\max_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))$$

Instead of minimizing likelihood of discriminator being correct, now maximize likelihood of discriminator being wrong.

Same objective of fooling discriminator, but now higher gradient signal for bad samples => works much better! Standard in practice.



x-axis:  $D(G(z))$ , y-axis (blue):  $\log(1 - D(G(z)))$ , (green):  $-\log(D(G(z)))$

# Training GANs: Two-player game

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. Gradient ascent on discriminator

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

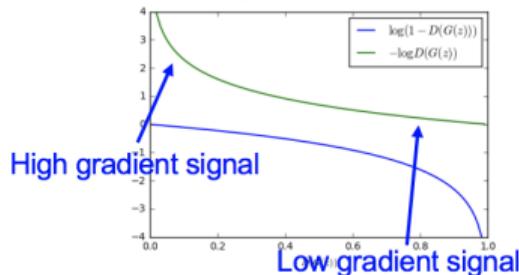
2. Instead: Gradient ascent on generator, different objective

$$\max_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))$$

Instead of minimizing likelihood of discriminator being correct, now maximize likelihood of discriminator being wrong.

Same objective of fooling discriminator, but now higher gradient signal for bad samples => works much better! Standard in practice.

Aside: Jointly training two networks is challenging, can be unstable. Choosing objectives with better loss landscapes helps training, is an active area of research.



x-axis:  $D(G(z))$ , y-axis (blue):  $\log(1 - D(G(z)))$ , (green):  $-\log(D(G(z)))$

# Training GANs: Two-player game

## Putting it together: GAN training algorithm

```
for number of training iterations do
    for k steps do
        • Sample minibatch of m noise samples { $\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}$ } from noise prior  $p_g(\mathbf{z})$ .
        • Sample minibatch of m examples { $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}$ } from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
        • Update the discriminator by ascending its stochastic gradient:
```

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D_{\theta_d}(x^{(i)}) + \log(1 - D_{\theta_d}(G_{\theta_g}(z^{(i)}))) \right]$$

```
    end for
    • Sample minibatch of m noise samples { $\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}$ } from noise prior  $p_g(\mathbf{z})$ .
    • Update the generator by ascending its stochastic gradient (improved objective):
```

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(D_{\theta_d}(G_{\theta_g}(z^{(i)})))$$

```
end for
```

# Training GANs: Two-player game

## Putting it together: GAN training algorithm

for number of training iterations do

for  $k$  steps do

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D_{\theta_d}(\mathbf{x}^{(i)}) + \log(1 - D_{\theta_d}(G_{\theta_g}(\mathbf{z}^{(i)}))) \right]$$

Some find  $k=1$   
more stable,  
others use  $k > 1$ ,  
no best rule.

Recent work (e.g.  
Wasserstein GAN)  
alleviates this  
problem, better  
stability!

end for

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Update the generator by ascending its stochastic gradient (improved objective):

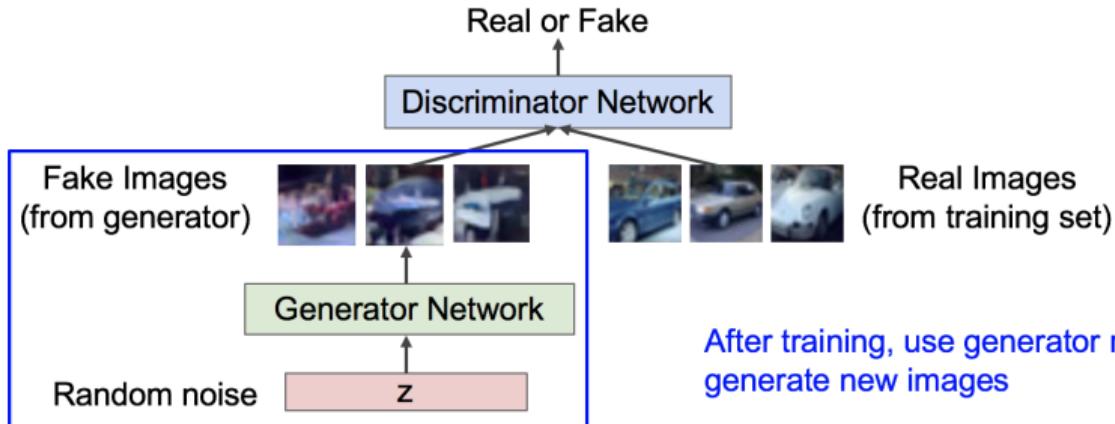
$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(D_{\theta_d}(G_{\theta_g}(\mathbf{z}^{(i)})))$$

end for

## Training GANs: Two-player game

**Generator network:** try to fool the discriminator by generating real-looking images

**Discriminator network:** try to distinguish between real and fake images



Fake and real images copyright Emily Denton et al. 2015. Reproduced with permission.

## Jensen-Shannon Divergence

### Definition 1 (Jensen-Shannon Divergence)

Jensen-Shannon Divergence (JSD) is defined by

$$JSD(P \parallel Q) = \frac{1}{2}KL(P \parallel M) + \frac{1}{2}KL(Q \parallel M), \text{ where } M = \frac{1}{2}(P + Q)$$

- ▶ Note, the JSD is bounded below by zero and above by 1 (base 2) or  $\ln(2)$  (base  $e$ )
- ▶ Side note, KL is not bounded above and can reach  $\infty$

## Optimal discriminator $D$ for any given generator $G$

### Proposition 2

*For  $G$  fixed, the optimal discriminator  $D$  is*

$$D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_g(x)} \quad (1)$$

## Proof.

The training criterion for the discriminator  $D$ , given any generator  $G$ , is to maximize the quantity  $V(G, D)$

$$\begin{aligned} V(G, D) &= \int_x p_{data}(x) \log(D(x)) dx + \int_z p_z(z) \log(1 - D(G(z))) dz \\ &= \int_x p_{data}(x) \log(D(x)) + p_g(x) \log(1 - D(x)) dx \end{aligned}$$

For any  $(a, b) \in \mathbb{R}^2 \setminus \{0, 0\}$ , the function  $y \rightarrow a \log(y) + b \log(1 - y)$  achieves its maximum in  $[0, 1]$  at  $\frac{a}{a+b}$ . The discriminator does not need to be defined outside of  $Supp(p_{data}) \cup Supp(p_g)$ . ■

- ▶ Note that the training objective for  $D$  can be interpreted as maximizing the log-likelihood for estimating the conditional probability  $P(Y = y | x)$ , where  $Y$  indicates whether  $x$  comes from  $p_{data}$  (with  $y = 1$ ) or from  $p_g$  (with  $y = 0$ ). The minimax game in Equation (1) can now be reformulated as following
- ▶

$$\begin{aligned}
 C(G) &= \max_D V(G, D) = \mathbb{E}_{x \sim p_{data}} \log D_G^*(x) + \mathbb{E}_{z \sim p_z} \log(1 - D_G^*(G(z))) \\
 &= \mathbb{E}_{x \sim p_{data}} \log D_G^*(x) + \mathbb{E}_{x \sim p_g} \log(1 - D_G^*(x)) \\
 &= \mathbb{E}_{x \sim p_{data}} \log \frac{p_{data}(x)}{p_{data}(x) + p_g(x)} + \mathbb{E}_{x \sim p_g} \log \frac{p_g(x)}{p_{data}(x) + p_g(x)} \\
 &= -\log(4) + KL \left( p_{data} \parallel \frac{p_{data} + p_g}{2} \right) + KL \left( p_g \parallel \frac{p_{data} + p_g}{2} \right) \\
 &= -\log(4) + 2JSD(p_{data} \parallel p_g)
 \end{aligned} \tag{2}$$

## Global optimality of $p_g = p_{data}$

### Theorem 3

*The global minimum of the training criterion  $C(G)$  in Equation (2) is achieved iff  $p_g = p_{data}$ . At that point,  $C(G)$  achieves the value of  $-\log(4)$ .*

## Proof.

For  $p_g = p_{data}$ , it is obvious  $D_G^*(x) = 1/2$  from Equation (1). Hence, by inspecting Equation (2) at  $D_G^*(x) = 1/2$ , we find

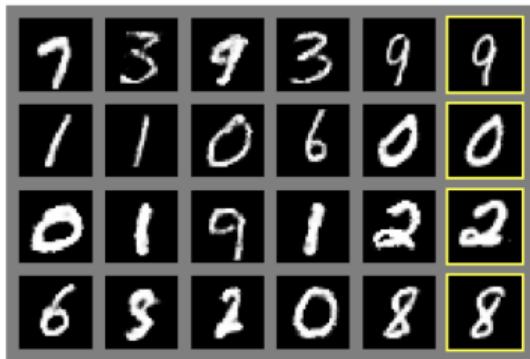
$C(G) = \log \frac{1}{2} + \log \frac{1}{2} = -\log(4)$ . To see that this is the best possible value of  $C(G)$ , reached only for  $p_g = p_{data}$ , observe that

$$\mathbb{E}_{x \sim p_{data}} - \log 2 + \mathbb{E}_{x \sim p_g} - \log 2 = -\log 4$$

and since Jensen-Shannon divergence between two distributions is always non-negative, and zero iff they are equal, we conclude from Equation (2) that  $C^* = -\log(4)$  is the global minimum of  $C(G)$  and that the only solution is  $p_g = p_{data}$ , i.e. the generative model perfectly replicating the data distribution. ■

# Generative Adversarial Nets

Generated samples



Nearest neighbor from training set



Figures copyright Ian Goodfellow et al., 2014. Reproduced with permission.

# Generative Adversarial Nets

Generated samples (CIFAR-10)



Nearest neighbor from training set

Figures copyright Ian Goodfellow et al., 2014. Reproduced with permission.

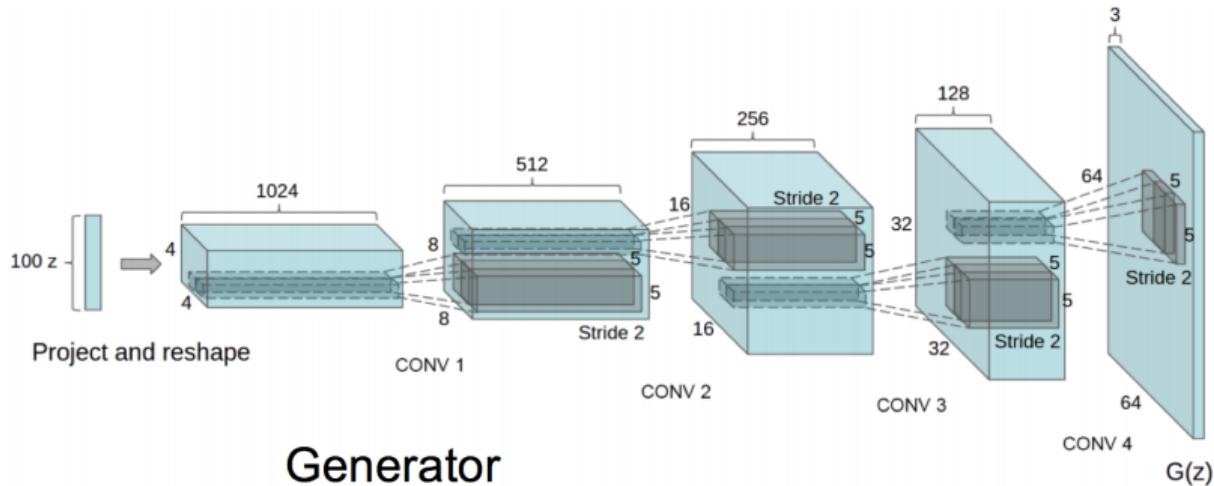
# DCGAN

Architecture guidelines for stable Deep Convolutional GANs

- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses Tanh.
- Use LeakyReLU activation in the discriminator for all layers.

Radford et al, "Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks", ICLR 2016

# DCGAN

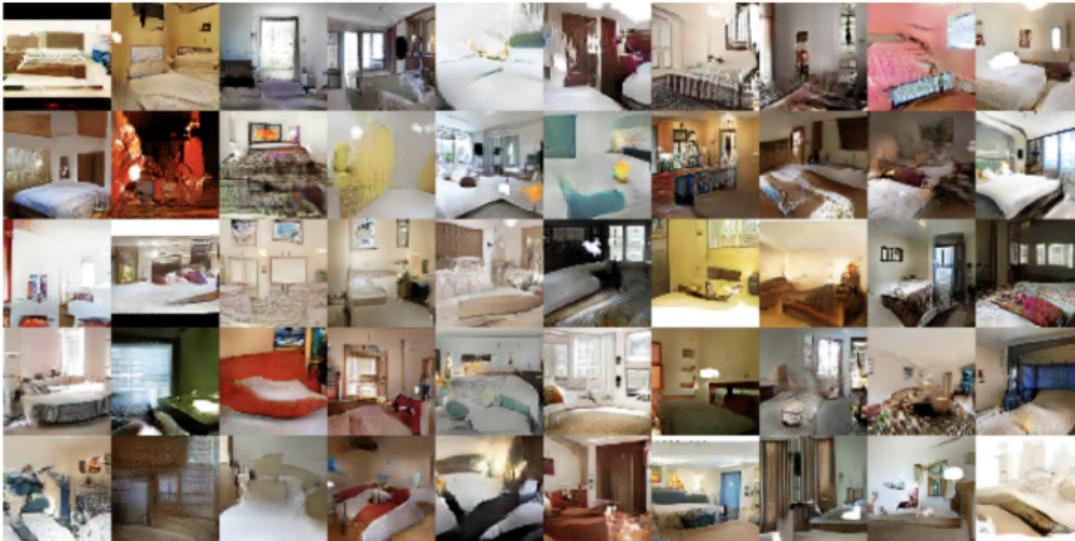


Radford et al, "Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks", ICLR 2016

# DCGAN

Samples  
from the  
model look  
amazing!

Radford et al,  
ICLR 2016



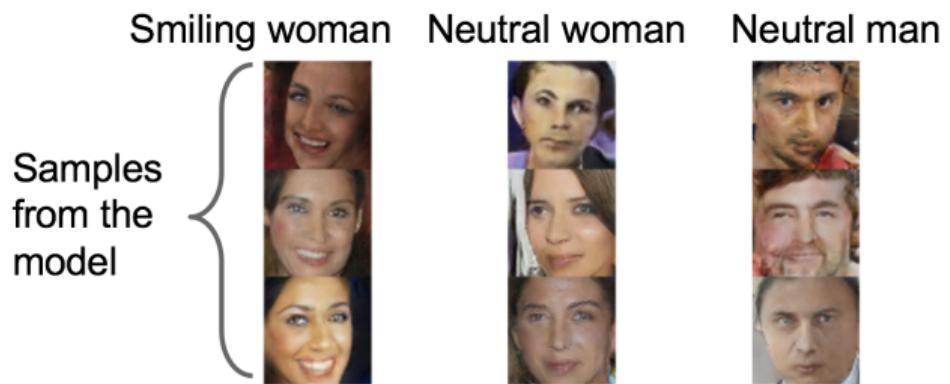
# DCGAN

Interpolating  
between  
random  
points in  
latent space

Radford et al,  
ICLR 2016

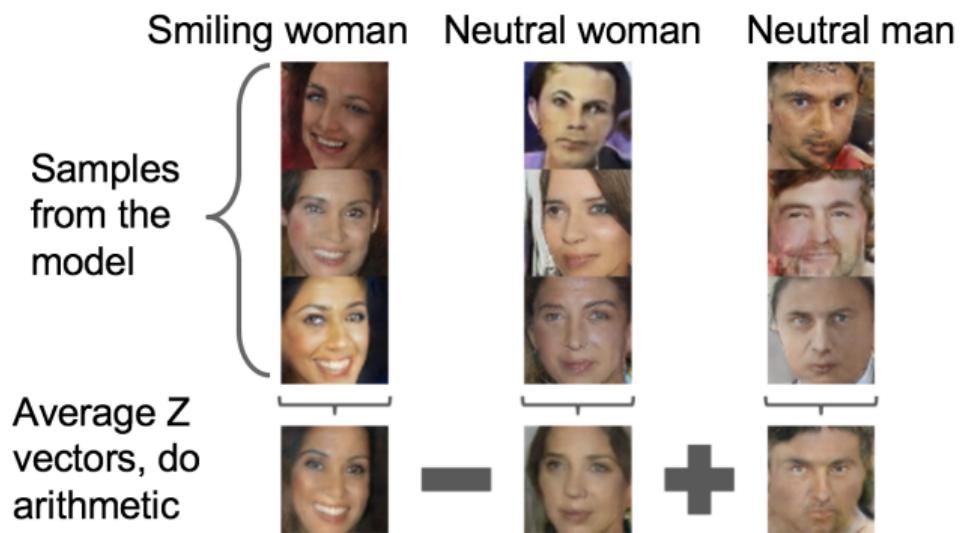


## GAN vector math



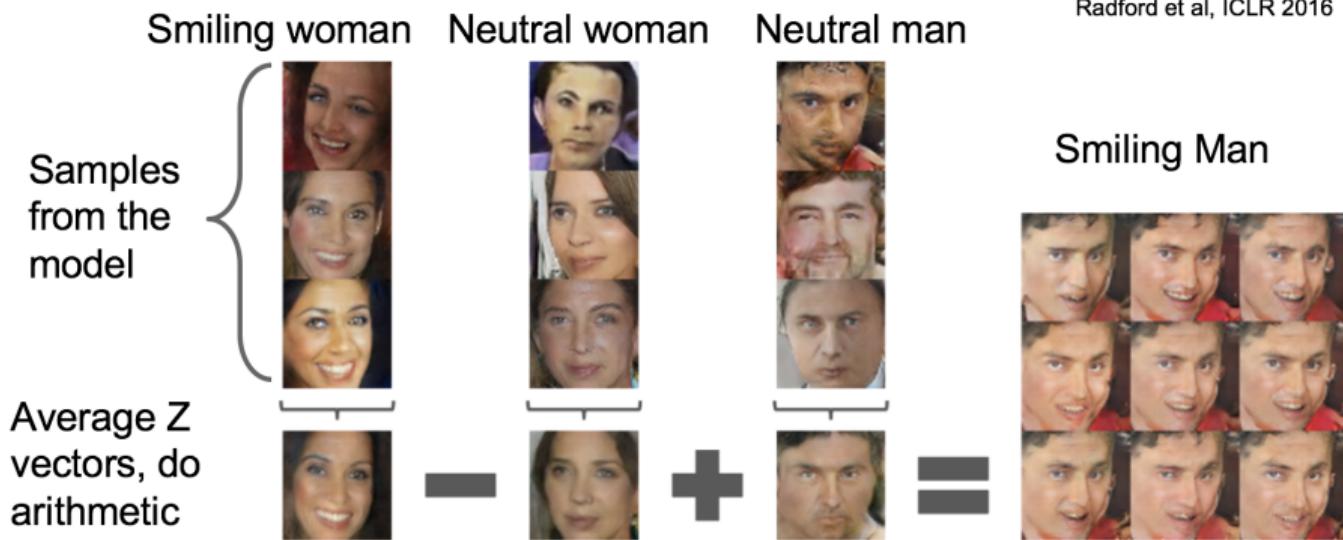
Radford et al, ICLR 2016

## GAN vector math



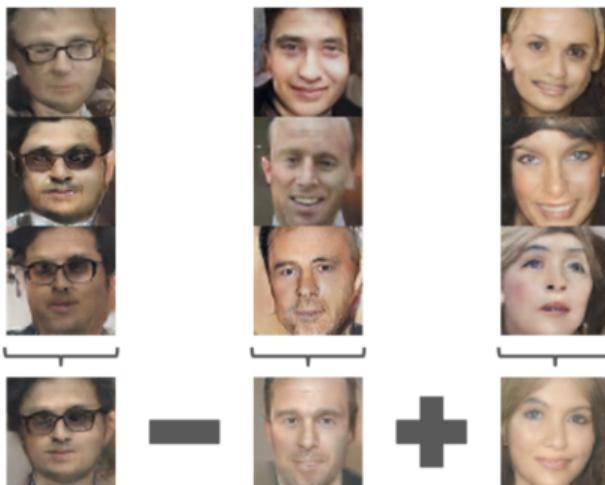
Radford et al, ICLR 2016

## GAN vector math



## GAN vector math

Glasses man    No glasses man    No glasses woman



Radford et al,  
ICLR 2016

## GAN vector math

Glasses man



No glasses man

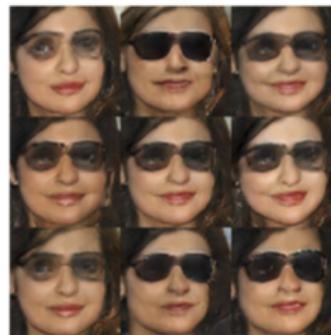


No glasses woman



Radford et al,  
ICLR 2016

Woman with glasses



# “The GAN Zoo”

- GAN - Generative Adversarial Networks
- 3D-GAN - Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling
- acGAN - Face Aging With Conditional Generative Adversarial Networks
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- AdaGAN - AdaGAN: Boosting Generative Models
- AEGAN - Learning Inverse Mapping by Autoencoder based Generative Adversarial Nets
- AffGAN - Amortized MAP Inference for Image Super-resolution
- AL-CGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference
- AM-GAN - Generative Adversarial Nets with Labeled Data by Activation Maximization
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- ArtGAN - ArtGAN: Artwork Synthesis with Conditional Categorical GANs
- b-GAN - b-GAN: Unified Framework of Generative Adversarial Networks
- Bayesian GAN - Deep and Hierarchical Implicit Models
- BEGAN - BEGAN: Boundary Equilibrium Generative Adversarial Networks
- BiGAN - Adversarial Feature Learning
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- CGAN - Conditional Generative Adversarial Nets
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks
- CCGAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CoGAN - Coupled Generative Adversarial Networks
- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
- DTN - Unsupervised Cross-Domain Image Generation
- DCGAN - Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
- DiscoGAN - Learning to Discover Cross-Domain Relations with Generative Adversarial Networks
- DR-GAN - Disentangled Representation Learning GAN for Pose-Invariant Face Recognition
- DualGAN - DualGAN: Unsupervised Dual Learning for Image-to-Image Translation
- EBGAN - Energy-based Generative Adversarial Network
- f-GAN - f-GAN: Training Generative Neural Samplers using Variational Divergence Minimization
- FF-GAN - Towards Large-Pose Face Frontalization in the Wild
- GAWWN - Learning What and Where to Draw
- GeneGAN - GeneGAN: Learning Object Transfiguration and Attribute Subspace from Unpaired Data
- Geometric GAN - Geometric GAN
- GoGAN - Gang of GANs: Generative Adversarial Networks with Maximum Margin Ranking
- GP-GAN - GP-GAN: Towards Realistic High-Resolution Image Blending
- IAN - Neural Photo Editing with Introspective Adversarial Networks
- iGAN - Generative Visual Manipulation on the Natural Image Manifold
- IcGAN - Invertible Conditional GANs for image editing
- ID-CGAN - Image De-raining Using a Conditional Generative Adversarial Network
- Improved GAN - Improved Techniques for Training GANs
- InfoGAN - InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets
- LAGAN - Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis
- LAPGAN - Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks

See also: <https://github.com/soumith/ganhacks> for tips and tricks for trainings GANs

<https://github.com/hindupuravinash/the-gan-zoo>

# Outline

Autoregressive models: PixelRNN and PixelCNN

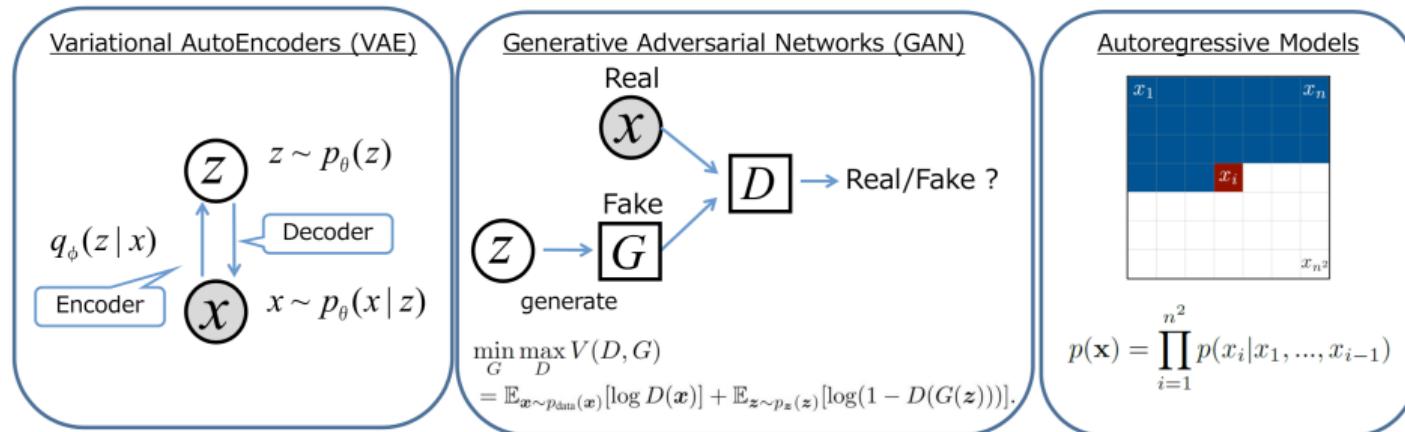
Generative adversarial networks

Summary of deep generative models

# Comparisons<sup>4</sup>

## Image Generation Models

-Three image generation approaches are dominating the field:



|       | VAE  | GAN   | Autoregressive Models   |
|-------|--|---|---|
| Pros. | - Efficient inference with approximate latent variables. | - generate sharp image.<br>- no need for any Markov chain or approx networks during sampling. | - very simple and stable training process<br>- currently gives the best log likelihood.<br>- tractable likelihood |
| Cons. | - generated samples tend to be blurry.                   | - difficult to optimize due to unstable training dynamics.                                    | - relatively inefficient during sampling  |

<sup>4</sup> from Yohei Sugawara

(cf. <https://openai.com/blog/generative-models/>)

## Comparisons<sup>5</sup>

- ▶ **Variational Autoencoders** allow us to perform both learning and efficient Bayesian inference in sophisticated probabilistic graphical models with latent variables. However, their generated samples tend to be slightly blurry.

---

<sup>5</sup><https://openai.com/blog/generative-models/>  
Summary of deep generative models

## Comparisons<sup>5</sup>

- ▶ **Variational Autoencoders** allow us to perform both learning and efficient Bayesian inference in sophisticated probabilistic graphical models with latent variables. However, their generated samples tend to be slightly blurry.
- ▶ **GANs** currently generate the sharpest images but they are more difficult to optimize due to unstable training dynamics.

---

<sup>5</sup><https://openai.com/blog/generative-models/>

## Comparisons<sup>5</sup>

- ▶ **Variational Autoencoders** allow us to perform both learning and efficient Bayesian inference in sophisticated probabilistic graphical models with latent variables. However, their generated samples tend to be slightly blurry.
- ▶ **GANs** currently generate the sharpest images but they are more difficult to optimize due to unstable training dynamics.
- ▶ **PixelRNNs** have a very simple and stable training process (softmax loss) and currently give the best log likelihoods (that is, plausibility of the generated data). However, they are relatively inefficient during sampling and don't easily provide simple low-dimensional codes for images.
- ▶ All of these models are active areas of research and we are eager to see how they develop in the future!

---

<sup>5</sup><https://openai.com/blog/generative-models/>