

VAE

Variational autoencoder [1] models inherit autoencoder architecture, but use variational approach
homework, we will implement VAE and quantitatively measure the quality of the generated sample

[1] Auto-Encoding Variational Bayes, Diederik P Kingma, Max Welling 2013 <https://arxiv.org/abs/1312.6114>

[2] Improved techniques for training gans, Salimans, T., Goodfellow, I., Zaremba, W., Cheung, V., Radford, A.
Neural Information Processing Systems

[3] A note on inception score, Shane Barratt, Rishi Sharma 2018 <https://arxiv.org/abs/1801.01973>

▼ PART I. Train a good VAE model

▼ Setup

```
import tensorflow as tf
if tf.__version__ < '2.0.0':
    tf.enable_eager_execution()
tf.executing_eagerly()

import numpy as np
import os

import matplotlib.pyplot as plt
import matplotlib.gridspec as gridspec

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# A bunch of utility functions

def show_images(images):
    # images reshape to (batch_size, D)
    images = np.reshape(images, [images.shape[0], -1])
    sqrt_n = int(np.ceil(np.sqrt(images.shape[0])))
    sqrt_m = int(np.ceil(np.sqrt(images.shape[1])))

    fig = plt.figure(figsize=(sqrt_n, sqrt_m))
    gs = gridspec.GridSpec(sqrt_n, sqrt_m)
    gs.update(wspace=0.05, hspace=0.05)

    for i, img in enumerate(images):
        ax = plt.subplot(gs[i])
        plt.axis('off')
        ax.set_xticklabels([])
```

```

        ax.set_yticklabels([])
        ax.set_aspect('equal')
        plt.imshow(img.reshape([sqrt(img), sqrt(img)]))
    return

def preprocess_img(x):
    return 2 * x - 1.0

def rel_error(x,y):
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

def count_params(model):
    """Count the number of parameters in the current TensorFlow graph """
    param_count = np.sum([np.prod(p.shape) for p in model.weights])
    return param_count

```

▼ Dataset

We will be working on the MNIST dataset, which is 60,000 training and 10,000 test images. Each pixel is on a black background (0 through 9). This was one of the first datasets used to train convolutional standard CNN model can easily exceed 99% accuracy.

Heads-up: Our MNIST wrapper returns images as vectors. That is, they're size (batch, 784). If you want to resize them to (batch,28,28) or (batch,28,28,1). They are also type np.float32 and bounded [0,1].

```

class MNIST(object):
    def __init__(self, batch_size, shuffle=False):
        """
        Construct an iterator object over the MNIST data

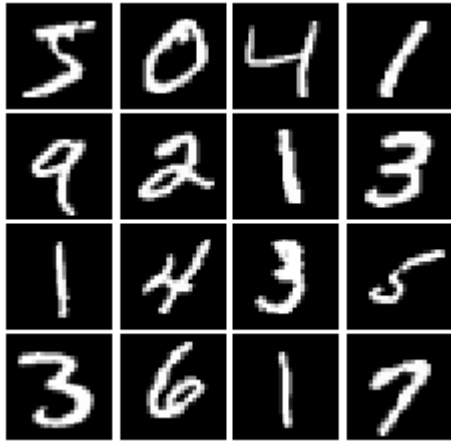
        Inputs:
        - batch_size: Integer giving number of elements per minibatch
        - shuffle: (optional) Boolean, whether to shuffle the data on each epoch
        """
        train, _ = tf.keras.datasets.mnist.load_data()
        X, y = train
        X = X.astype(np.float32)/255
        X = X.reshape((X.shape[0], -1))
        self.X, self.y = X, y
        self.batch_size, self.shuffle = batch_size, shuffle

    def __iter__(self):
        N, B = self.X.shape[0], self.batch_size
        idxs = np.arange(N)
        if self.shuffle:
            np.random.shuffle(idxs)
        return iter((self.X[i:i+B], self.y[i:i+B]) for i in range(0, N, B))

# show a batch
mnist = MNIST(batch_size=16)
show_images(mnist.X[:16])

```

↳ Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step



```
X_DIM = mnist.X[0].size
num_samples = 100000
num_to_show = 100

# Hyperparameters. Your job to find these.
# TODO:
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
num_epochs = 100
batch_size = 50
Z_DIM = 5
learning_rate = 0.0005
# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
```

▼ Encoder

Our first step is to build a variational encoder network $q_{\phi}(z | x)$.

Hint: You should use the layers in `tf.keras.layers` to build the model. Use four FC layers. All fully c
For initialization, just use the default initializer used by the `tf.keras.layers` functions.

The output of the encoder should thus have shape `[batch_size, 2*z_dim]`, and contain real number diagonal log variance $\log \sigma(x_i)^2$ of each of the `batch_size` input images. Note, we want to make i stability.

WARNING: Do not apply any non-linearity to the last activation.

```
def q_phi(z_dim=Z_DIM, x_dim=X_DIM):
    model = tf.keras.models.Sequential([
        # TODO: implement architecture
        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

        tf.keras.layers.Dense(400, input_shape = (x_dim, ), activation = 'relu'),
        tf.keras.layers.Dense(200, activation = 'relu'),
        tf.keras.layers.Dense(100, activation = 'tanh'),
        tf.keras.layers.Dense(2 * z_dim),

        # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    ])
```

```

    ..
    return model

# TODO: implement reparameterization trick
def sample_z(mu, log_var):
    # Your code here for the reparameterization trick.
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    samples = None

    z = tf.random.normal(tf.shape(mu))
    samples = mu + z * tf.math.exp(0.5*log_var)

    return samples
    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

```

▼ Decoder

Now to build a decoder network $p_{\theta}(x | z)$. You should use the layers in `tf.keras.layers` to construct connected layers should include bias terms. Note that you can use the `tf.nn` module to access activation functions and initializers for parameters.

In this exercise, we will use Bernoulli MLP decoder where $p_{\theta}(x | z)$ is modeled with multivariate Bernoulli distribution we discussed in the lecture, as following (see Appendix C.1 in the original paper)

$$\log p(x | z) = \sum_{i=1} x_i \log z_i + (1 - x_i) \log(1 - z_i)$$

Note, the output of the decoder should have shape `[batch_size, x_dim]` and should output the unnormalized probabilities.

WARNING: Do not apply any non-linearity to the last activation.

```

def p_theta(z_dim=Z_DIM, x_dim=X_DIM):
    model = tf.keras.models.Sequential([
        # TODO: implement architecture
        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

        tf.keras.layers.Dense(100, input_shape = (z_dim, ), activation = 'tanh'),
        tf.keras.layers.Dense(200, activation = 'relu'),
        tf.keras.layers.Dense(400, activation = 'relu'),
        tf.keras.layers.Dense(x_dim),

        # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    ])
    return model

```

▼ Loss definition

Compute the VAE loss.

1. For the reconstruction loss, you might find `tf.nn.sigmoid_cross_entropy_with_logits` or `tf.keras.losses.binary_crossentropy`.
2. For the kl loss, we discussed the closed form kl divergence between two gaussians in the lecture.

```
def vae_loss(x, x_logit, z_mu, z_logvar):
    recon_loss = None
    kl_loss = None

    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

    recon_loss = tf.reduce_sum(tf.nn.sigmoid_cross_entropy_with_logits(x, x_logit), axis = 1)
    kl_loss = -0.5 * tf.reduce_sum(1 + z_logvar - tf.square(z_mu) - tf.math.exp(z_logvar), axis = 1)

    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    vae_loss = tf.reduce_mean(recon_loss + kl_loss)
    return vae_loss, tf.reduce_mean(recon_loss)
```


▼ Optimizing our loss

```
Q = q_phi()
P = p_theta()
solver = tf.keras.optimizers.Adam(learning_rate)
mnist = MNIST(batch_size=batch_size, shuffle=True)
```

Visualize generated samples before training

```
z_gen = tf.random.normal(shape=[num_to_show, Z_DIM])
x_gen = P(z_gen)
imgs_numpy = tf.nn.sigmoid(x_gen).numpy()
show_images(imgs_numpy)
plt.show()
```





▼ Training a VAE!

If everything works, your batch average reconstruction loss should drop below 95.



```

iter_count = 0
show_every = 200
for epoch in range(num_epochs):
    for (x_i, _) in mnist:
        with tf.GradientTape() as tape:
            z_concat = Q(preprocess_img(x_i))
            z_mu, z_logvar = tf.split(z_concat, num_or_size_splits=2, axis=1)
            z_i = sample_z(z_mu, z_logvar)

            x_logit = P(z_i)

            loss, recon_loss = vae_loss(x_i, x_logit, z_mu, z_logvar)

            grads = tape.gradient(loss,
                                  [Q.trainable_variables, P.trainable_variables])

            solver.apply_gradients(zip([*grads[0], *grads[1]],
                                      [*Q.trainable_variables, *P.trainable_variables]))

        if (iter_count % show_every == 0):
            print('Epoch: {}, Iter: {}, Loss: {:.4}, Recon: {:.4}'.format(
                epoch, iter_count, loss, recon_loss))
            #imgs_numpy = tf.nn.sigmoid(x_logit).numpy()
            #show_images(imgs_numpy[0:16])
            #plt.show()
        iter_count += 1

```



Epoch: 0, Iter: 0, Loss: 546.6, Recon: 543.1
Epoch: 0, Iter: 200, Loss: 185.7, Recon: 179.4
Epoch: 0, Iter: 400, Loss: 172.2, Recon: 164.0
Epoch: 0, Iter: 600, Loss: 154.3, Recon: 145.5
Epoch: 0, Iter: 800, Loss: 144.3, Recon: 134.9
Epoch: 0, Iter: 1000, Loss: 136.1, Recon: 126.9
Epoch: 1, Iter: 1200, Loss: 136.0, Recon: 126.4
Epoch: 1, Iter: 1400, Loss: 137.9, Recon: 127.9
Epoch: 1, Iter: 1600, Loss: 144.9, Recon: 134.5
Epoch: 1, Iter: 1800, Loss: 138.0, Recon: 127.9
Epoch: 1, Iter: 2000, Loss: 135.2, Recon: 124.9
Epoch: 1, Iter: 2200, Loss: 126.4, Recon: 116.5
Epoch: 2, Iter: 2400, Loss: 128.2, Recon: 117.7
Epoch: 2, Iter: 2600, Loss: 132.2, Recon: 121.6
Epoch: 2, Iter: 2800, Loss: 136.3, Recon: 125.2
Epoch: 2, Iter: 3000, Loss: 134.0, Recon: 123.1
Epoch: 2, Iter: 3200, Loss: 131.7, Recon: 121.0
Epoch: 2, Iter: 3400, Loss: 123.2, Recon: 112.9
Epoch: 3, Iter: 3600, Loss: 123.6, Recon: 113.0
Epoch: 3, Iter: 3800, Loss: 129.6, Recon: 118.2
Epoch: 3, Iter: 4000, Loss: 131.6, Recon: 120.3
Epoch: 3, Iter: 4200, Loss: 129.6, Recon: 118.3
Epoch: 3, Iter: 4400, Loss: 128.5, Recon: 117.2
Epoch: 3, Iter: 4600, Loss: 118.9, Recon: 108.1
Epoch: 4, Iter: 4800, Loss: 120.8, Recon: 109.8
Epoch: 4, Iter: 5000, Loss: 128.3, Recon: 117.0
Epoch: 4, Iter: 5200, Loss: 128.1, Recon: 116.6
Epoch: 4, Iter: 5400, Loss: 128.0, Recon: 116.5
Epoch: 4, Iter: 5600, Loss: 125.0, Recon: 113.8
Epoch: 4, Iter: 5800, Loss: 116.9, Recon: 106.1
Epoch: 5, Iter: 6000, Loss: 118.5, Recon: 107.3
Epoch: 5, Iter: 6200, Loss: 125.3, Recon: 114.0
Epoch: 5, Iter: 6400, Loss: 126.4, Recon: 114.7
Epoch: 5, Iter: 6600, Loss: 126.7, Recon: 115.1
Epoch: 5, Iter: 6800, Loss: 122.7, Recon: 111.2
Epoch: 5, Iter: 7000, Loss: 115.6, Recon: 104.5
Epoch: 6, Iter: 7200, Loss: 117.5, Recon: 106.3
Epoch: 6, Iter: 7400, Loss: 124.6, Recon: 113.0
Epoch: 6, Iter: 7600, Loss: 125.3, Recon: 113.7
Epoch: 6, Iter: 7800, Loss: 126.0, Recon: 114.2
Epoch: 6, Iter: 8000, Loss: 122.2, Recon: 110.6
Epoch: 6, Iter: 8200, Loss: 115.0, Recon: 103.6
Epoch: 7, Iter: 8400, Loss: 115.6, Recon: 104.2
Epoch: 7, Iter: 8600, Loss: 124.6, Recon: 112.9
Epoch: 7, Iter: 8800, Loss: 123.4, Recon: 111.5
Epoch: 7, Iter: 9000, Loss: 124.1, Recon: 112.1
Epoch: 7, Iter: 9200, Loss: 120.7, Recon: 109.0
Epoch: 7, Iter: 9400, Loss: 113.1, Recon: 101.7
Epoch: 8, Iter: 9600, Loss: 114.4, Recon: 102.7
Epoch: 8, Iter: 9800, Loss: 123.0, Recon: 111.2
Epoch: 8, Iter: 10000, Loss: 122.6, Recon: 110.6
Epoch: 8, Iter: 10200, Loss: 123.8, Recon: 111.8
Epoch: 8, Iter: 10400, Loss: 119.7, Recon: 107.8
Epoch: 8, Iter: 10600, Loss: 113.0, Recon: 101.4
Epoch: 9, Iter: 10800, Loss: 114.4, Recon: 102.8
Epoch: 9, Iter: 11000, Loss: 122.0, Recon: 110.2
Epoch: 9, Iter: 11200, Loss: 122.6, Recon: 110.6
Epoch: 9, Iter: 11400, Loss: 123.7, Recon: 111.6

```
Epoch: 9, Iter: 11600, Loss: 118.3, Recon: 106.5
Epoch: 9, Iter: 11800, Loss: 111.9, Recon: 100.2
Epoch: 10, Iter: 12000, Loss: 112.2, Recon: 100.5
Epoch: 10, Iter: 12200, Loss: 122.4, Recon: 110.5
Epoch: 10, Iter: 12400, Loss: 120.8, Recon: 108.5
Epoch: 10, Iter: 12600, Loss: 122.7, Recon: 110.6
Epoch: 10, Iter: 12800, Loss: 118.2, Recon: 106.1
Epoch: 10, Iter: 13000, Loss: 110.6, Recon: 98.64
Epoch: 11, Iter: 13200, Loss: 111.0, Recon: 99.01
Epoch: 11, Iter: 13400, Loss: 121.0, Recon: 109.3
Epoch: 11, Iter: 13600, Loss: 121.1, Recon: 108.8
Epoch: 11, Iter: 13800, Loss: 121.5, Recon: 109.2
Epoch: 11, Iter: 14000, Loss: 118.8, Recon: 106.7
Epoch: 11, Iter: 14200, Loss: 110.8, Recon: 98.8
Epoch: 12, Iter: 14400, Loss: 112.3, Recon: 100.3
Epoch: 12, Iter: 14600, Loss: 120.7, Recon: 108.7
Epoch: 12, Iter: 14800, Loss: 119.0, Recon: 106.7
Epoch: 12, Iter: 15000, Loss: 120.5, Recon: 108.3
Epoch: 12, Iter: 15200, Loss: 116.8, Recon: 104.8
Epoch: 12, Iter: 15400, Loss: 111.0, Recon: 99.06
Epoch: 13, Iter: 15600, Loss: 110.9, Recon: 98.91
Epoch: 13, Iter: 15800, Loss: 120.9, Recon: 108.8
Epoch: 13, Iter: 16000, Loss: 118.9, Recon: 106.7
Epoch: 13, Iter: 16200, Loss: 120.7, Recon: 108.3
Epoch: 13, Iter: 16400, Loss: 116.4, Recon: 104.3
Epoch: 13, Iter: 16600, Loss: 110.9, Recon: 98.86
Epoch: 14, Iter: 16800, Loss: 110.8, Recon: 98.96
Epoch: 14, Iter: 17000, Loss: 120.3, Recon: 108.2
Epoch: 14, Iter: 17200, Loss: 119.1, Recon: 106.6
Epoch: 14, Iter: 17400, Loss: 119.7, Recon: 107.4
Epoch: 14, Iter: 17600, Loss: 115.4, Recon: 103.2
Epoch: 14, Iter: 17800, Loss: 110.3, Recon: 98.28
Epoch: 15, Iter: 18000, Loss: 109.3, Recon: 97.31
Epoch: 15, Iter: 18200, Loss: 120.1, Recon: 108.0
Epoch: 15, Iter: 18400, Loss: 117.5, Recon: 105.3
Epoch: 15, Iter: 18600, Loss: 119.0, Recon: 106.5
Epoch: 15, Iter: 18800, Loss: 114.7, Recon: 102.4
Epoch: 15, Iter: 19000, Loss: 109.7, Recon: 97.53
Epoch: 16, Iter: 19200, Loss: 109.3, Recon: 97.44
Epoch: 16, Iter: 19400, Loss: 119.8, Recon: 107.7
Epoch: 16, Iter: 19600, Loss: 117.3, Recon: 105.0
Epoch: 16, Iter: 19800, Loss: 118.8, Recon: 106.3
Epoch: 16, Iter: 20000, Loss: 115.6, Recon: 103.2
Epoch: 16, Iter: 20200, Loss: 109.5, Recon: 97.32
Epoch: 17, Iter: 20400, Loss: 109.9, Recon: 97.78
Epoch: 17, Iter: 20600, Loss: 119.3, Recon: 107.1
Epoch: 17, Iter: 20800, Loss: 117.2, Recon: 104.8
Epoch: 17, Iter: 21000, Loss: 119.0, Recon: 106.6
Epoch: 17, Iter: 21200, Loss: 116.1, Recon: 103.8
Epoch: 17, Iter: 21400, Loss: 109.8, Recon: 97.64
Epoch: 18, Iter: 21600, Loss: 112.6, Recon: 100.5
Epoch: 18, Iter: 21800, Loss: 119.7, Recon: 107.6
Epoch: 18, Iter: 22000, Loss: 116.5, Recon: 103.9
Epoch: 18, Iter: 22200, Loss: 118.2, Recon: 105.7
Epoch: 18, Iter: 22400, Loss: 114.5, Recon: 102.2
Epoch: 18, Iter: 22600, Loss: 109.6, Recon: 97.65
Epoch: 19, Iter: 22800, Loss: 109.8, Recon: 97.61
Epoch: 19, Iter: 23000, Loss: 119.3, Recon: 107.1
```


Epoch: 19, Iter: 23200, Loss: 117.1, Recon: 104.6
Epoch: 19, Iter: 23400, Loss: 118.2, Recon: 105.5
Epoch: 19, Iter: 23600, Loss: 114.2, Recon: 102.0
Epoch: 19, Iter: 23800, Loss: 108.2, Recon: 95.95
Epoch: 20, Iter: 24000, Loss: 109.6, Recon: 97.52
Epoch: 20, Iter: 24200, Loss: 119.1, Recon: 106.9
Epoch: 20, Iter: 24400, Loss: 117.0, Recon: 104.6
Epoch: 20, Iter: 24600, Loss: 118.3, Recon: 105.6
Epoch: 20, Iter: 24800, Loss: 114.9, Recon: 102.7
Epoch: 20, Iter: 25000, Loss: 109.0, Recon: 96.93
Epoch: 21, Iter: 25200, Loss: 108.2, Recon: 96.04
Epoch: 21, Iter: 25400, Loss: 118.2, Recon: 106.1
Epoch: 21, Iter: 25600, Loss: 116.3, Recon: 103.8
Epoch: 21, Iter: 25800, Loss: 117.8, Recon: 105.2
Epoch: 21, Iter: 26000, Loss: 115.5, Recon: 103.2
Epoch: 21, Iter: 26200, Loss: 109.0, Recon: 96.67
Epoch: 22, Iter: 26400, Loss: 108.1, Recon: 96.06
Epoch: 22, Iter: 26600, Loss: 117.4, Recon: 105.1
Epoch: 22, Iter: 26800, Loss: 115.7, Recon: 103.1
Epoch: 22, Iter: 27000, Loss: 117.9, Recon: 105.3
Epoch: 22, Iter: 27200, Loss: 113.1, Recon: 100.9
Epoch: 22, Iter: 27400, Loss: 108.0, Recon: 95.59
Epoch: 23, Iter: 27600, Loss: 107.8, Recon: 95.61
Epoch: 23, Iter: 27800, Loss: 117.3, Recon: 105.1
Epoch: 23, Iter: 28000, Loss: 116.8, Recon: 103.9
Epoch: 23, Iter: 28200, Loss: 116.4, Recon: 103.9
Epoch: 23, Iter: 28400, Loss: 113.7, Recon: 101.3
Epoch: 23, Iter: 28600, Loss: 108.7, Recon: 96.44
Epoch: 24, Iter: 28800, Loss: 109.6, Recon: 97.52
Epoch: 24, Iter: 29000, Loss: 119.0, Recon: 106.8
Epoch: 24, Iter: 29200, Loss: 114.7, Recon: 102.1
Epoch: 24, Iter: 29400, Loss: 117.5, Recon: 104.9
Epoch: 24, Iter: 29600, Loss: 114.1, Recon: 101.5
Epoch: 24, Iter: 29800, Loss: 108.7, Recon: 96.34
Epoch: 25, Iter: 30000, Loss: 107.2, Recon: 94.89
Epoch: 25, Iter: 30200, Loss: 117.3, Recon: 105.1
Epoch: 25, Iter: 30400, Loss: 116.1, Recon: 103.5
Epoch: 25, Iter: 30600, Loss: 116.5, Recon: 103.7
Epoch: 25, Iter: 30800, Loss: 114.8, Recon: 102.4
Epoch: 25, Iter: 31000, Loss: 108.2, Recon: 95.73
Epoch: 26, Iter: 31200, Loss: 107.2, Recon: 94.83
Epoch: 26, Iter: 31400, Loss: 117.9, Recon: 105.4
Epoch: 26, Iter: 31600, Loss: 114.6, Recon: 101.9
Epoch: 26, Iter: 31800, Loss: 115.5, Recon: 102.8
Epoch: 26, Iter: 32000, Loss: 114.0, Recon: 101.5
Epoch: 26, Iter: 32200, Loss: 108.1, Recon: 95.65
Epoch: 27, Iter: 32400, Loss: 107.1, Recon: 94.82
Epoch: 27, Iter: 32600, Loss: 116.8, Recon: 104.4
Epoch: 27, Iter: 32800, Loss: 116.0, Recon: 103.1
Epoch: 27, Iter: 33000, Loss: 116.8, Recon: 104.1
Epoch: 27, Iter: 33200, Loss: 114.2, Recon: 101.8
Epoch: 27, Iter: 33400, Loss: 106.9, Recon: 94.52
Epoch: 28, Iter: 33600, Loss: 106.2, Recon: 93.92
Epoch: 28, Iter: 33800, Loss: 117.5, Recon: 105.0
Epoch: 28, Iter: 34000, Loss: 114.7, Recon: 101.8
Epoch: 28, Iter: 34200, Loss: 116.8, Recon: 104.1
Epoch: 28, Iter: 34400, Loss: 113.2, Recon: 100.7
Epoch: 28, Iter: 34600, Loss: 107.5, Recon: 95.18

Epoch: 29, Iter: 34800, Loss: 106.7, Recon: 94.47
Epoch: 29, Iter: 35000, Loss: 117.6, Recon: 105.1
Epoch: 29, Iter: 35200, Loss: 114.9, Recon: 102.1
Epoch: 29, Iter: 35400, Loss: 117.0, Recon: 104.4
Epoch: 29, Iter: 35600, Loss: 112.2, Recon: 99.85
Epoch: 29, Iter: 35800, Loss: 108.1, Recon: 95.73
Epoch: 30, Iter: 36000, Loss: 107.6, Recon: 95.39
Epoch: 30, Iter: 36200, Loss: 118.5, Recon: 105.9
Epoch: 30, Iter: 36400, Loss: 112.9, Recon: 100.1
Epoch: 30, Iter: 36600, Loss: 116.9, Recon: 104.2
Epoch: 30, Iter: 36800, Loss: 113.1, Recon: 100.7
Epoch: 30, Iter: 37000, Loss: 107.5, Recon: 95.0
Epoch: 31, Iter: 37200, Loss: 107.2, Recon: 94.83
Epoch: 31, Iter: 37400, Loss: 117.3, Recon: 104.9
Epoch: 31, Iter: 37600, Loss: 114.4, Recon: 101.5
Epoch: 31, Iter: 37800, Loss: 116.2, Recon: 103.5
Epoch: 31, Iter: 38000, Loss: 112.4, Recon: 99.69
Epoch: 31, Iter: 38200, Loss: 107.7, Recon: 95.4
Epoch: 32, Iter: 38400, Loss: 107.3, Recon: 94.83
Epoch: 32, Iter: 38600, Loss: 116.6, Recon: 104.0
Epoch: 32, Iter: 38800, Loss: 113.7, Recon: 100.9
Epoch: 32, Iter: 39000, Loss: 116.5, Recon: 103.8
Epoch: 32, Iter: 39200, Loss: 112.1, Recon: 99.47
Epoch: 32, Iter: 39400, Loss: 107.1, Recon: 94.61
Epoch: 33, Iter: 39600, Loss: 106.5, Recon: 94.08
Epoch: 33, Iter: 39800, Loss: 115.9, Recon: 103.5
Epoch: 33, Iter: 40000, Loss: 114.4, Recon: 101.6
Epoch: 33, Iter: 40200, Loss: 114.3, Recon: 101.5
Epoch: 33, Iter: 40400, Loss: 112.4, Recon: 99.9
Epoch: 33, Iter: 40600, Loss: 107.4, Recon: 95.01
Epoch: 34, Iter: 40800, Loss: 106.3, Recon: 94.02
Epoch: 34, Iter: 41000, Loss: 117.3, Recon: 104.9
Epoch: 34, Iter: 41200, Loss: 115.1, Recon: 102.2
Epoch: 34, Iter: 41400, Loss: 116.0, Recon: 103.4
Epoch: 34, Iter: 41600, Loss: 112.8, Recon: 100.2
Epoch: 34, Iter: 41800, Loss: 107.3, Recon: 94.66
Epoch: 35, Iter: 42000, Loss: 106.1, Recon: 93.6
Epoch: 35, Iter: 42200, Loss: 116.3, Recon: 103.7
Epoch: 35, Iter: 42400, Loss: 114.1, Recon: 101.2
Epoch: 35, Iter: 42600, Loss: 115.5, Recon: 102.8
Epoch: 35, Iter: 42800, Loss: 112.4, Recon: 99.91
Epoch: 35, Iter: 43000, Loss: 107.1, Recon: 94.51
Epoch: 36, Iter: 43200, Loss: 106.7, Recon: 94.35
Epoch: 36, Iter: 43400, Loss: 115.8, Recon: 103.2
Epoch: 36, Iter: 43600, Loss: 114.4, Recon: 101.6
Epoch: 36, Iter: 43800, Loss: 115.4, Recon: 102.6
Epoch: 36, Iter: 44000, Loss: 112.2, Recon: 99.69
Epoch: 36, Iter: 44200, Loss: 108.6, Recon: 95.98
Epoch: 37, Iter: 44400, Loss: 106.7, Recon: 94.2
Epoch: 37, Iter: 44600, Loss: 115.8, Recon: 103.3
Epoch: 37, Iter: 44800, Loss: 114.8, Recon: 101.9
Epoch: 37, Iter: 45000, Loss: 114.3, Recon: 101.4
Epoch: 37, Iter: 45200, Loss: 112.8, Recon: 100.2
Epoch: 37, Iter: 45400, Loss: 106.6, Recon: 94.11
Epoch: 38, Iter: 45600, Loss: 106.7, Recon: 94.39
Epoch: 38, Iter: 45800, Loss: 115.4, Recon: 102.8
Epoch: 38, Iter: 46000, Loss: 113.8, Recon: 100.7
Epoch: 38, Iter: 46200, Loss: 114.3, Recon: 101.4

Epoch: 38, Iter: 46400, Loss: 111.3, Recon: 98.61
Epoch: 38, Iter: 46600, Loss: 107.2, Recon: 94.67
Epoch: 39, Iter: 46800, Loss: 105.9, Recon: 93.3
Epoch: 39, Iter: 47000, Loss: 115.9, Recon: 103.4
Epoch: 39, Iter: 47200, Loss: 114.2, Recon: 101.2
Epoch: 39, Iter: 47400, Loss: 113.8, Recon: 101.0
Epoch: 39, Iter: 47600, Loss: 112.6, Recon: 100.0
Epoch: 39, Iter: 47800, Loss: 106.9, Recon: 94.34
Epoch: 40, Iter: 48000, Loss: 106.0, Recon: 93.45
Epoch: 40, Iter: 48200, Loss: 115.7, Recon: 103.0
Epoch: 40, Iter: 48400, Loss: 112.8, Recon: 99.9
Epoch: 40, Iter: 48600, Loss: 114.6, Recon: 101.7
Epoch: 40, Iter: 48800, Loss: 112.2, Recon: 99.51
Epoch: 40, Iter: 49000, Loss: 106.8, Recon: 94.27
Epoch: 41, Iter: 49200, Loss: 106.6, Recon: 94.07
Epoch: 41, Iter: 49400, Loss: 114.1, Recon: 101.4
Epoch: 41, Iter: 49600, Loss: 112.6, Recon: 99.57
Epoch: 41, Iter: 49800, Loss: 113.6, Recon: 100.7
Epoch: 41, Iter: 50000, Loss: 112.2, Recon: 99.57
Epoch: 41, Iter: 50200, Loss: 107.1, Recon: 94.42
Epoch: 42, Iter: 50400, Loss: 105.6, Recon: 93.19
Epoch: 42, Iter: 50600, Loss: 114.4, Recon: 101.9
Epoch: 42, Iter: 50800, Loss: 113.7, Recon: 100.7
Epoch: 42, Iter: 51000, Loss: 115.0, Recon: 102.2
Epoch: 42, Iter: 51200, Loss: 112.6, Recon: 100.1
Epoch: 42, Iter: 51400, Loss: 107.0, Recon: 94.33
Epoch: 43, Iter: 51600, Loss: 106.3, Recon: 93.74
Epoch: 43, Iter: 51800, Loss: 114.9, Recon: 102.3
Epoch: 43, Iter: 52000, Loss: 113.5, Recon: 100.5
Epoch: 43, Iter: 52200, Loss: 114.4, Recon: 101.5
Epoch: 43, Iter: 52400, Loss: 111.5, Recon: 99.04
Epoch: 43, Iter: 52600, Loss: 106.9, Recon: 94.37
Epoch: 44, Iter: 52800, Loss: 105.0, Recon: 92.52
Epoch: 44, Iter: 53000, Loss: 114.8, Recon: 102.1
Epoch: 44, Iter: 53200, Loss: 113.2, Recon: 100.2
Epoch: 44, Iter: 53400, Loss: 114.0, Recon: 101.1
Epoch: 44, Iter: 53600, Loss: 111.4, Recon: 98.86
Epoch: 44, Iter: 53800, Loss: 107.5, Recon: 94.84
Epoch: 45, Iter: 54000, Loss: 105.5, Recon: 93.06
Epoch: 45, Iter: 54200, Loss: 115.9, Recon: 103.3
Epoch: 45, Iter: 54400, Loss: 113.7, Recon: 100.8
Epoch: 45, Iter: 54600, Loss: 115.9, Recon: 103.2
Epoch: 45, Iter: 54800, Loss: 111.9, Recon: 99.37
Epoch: 45, Iter: 55000, Loss: 106.7, Recon: 93.96
Epoch: 46, Iter: 55200, Loss: 105.1, Recon: 92.57
Epoch: 46, Iter: 55400, Loss: 115.6, Recon: 103.1
Epoch: 46, Iter: 55600, Loss: 113.2, Recon: 100.0
Epoch: 46, Iter: 55800, Loss: 114.4, Recon: 101.5
Epoch: 46, Iter: 56000, Loss: 112.3, Recon: 99.76
Epoch: 46, Iter: 56200, Loss: 106.2, Recon: 93.49
Epoch: 47, Iter: 56400, Loss: 105.4, Recon: 92.88
Epoch: 47, Iter: 56600, Loss: 114.8, Recon: 102.1
Epoch: 47, Iter: 56800, Loss: 113.5, Recon: 100.5
Epoch: 47, Iter: 57000, Loss: 113.3, Recon: 100.3
Epoch: 47, Iter: 57200, Loss: 111.8, Recon: 99.24
Epoch: 47, Iter: 57400, Loss: 106.5, Recon: 93.83
Epoch: 48, Iter: 57600, Loss: 106.1, Recon: 93.57
Epoch: 48, Iter: 57800, Loss: 115.0, Recon: 102.5

Epoch: 48, Iter: 58000, Loss: 112.8, Recon: 99.79
Epoch: 48, Iter: 58200, Loss: 114.2, Recon: 101.2
Epoch: 48, Iter: 58400, Loss: 110.7, Recon: 98.04
Epoch: 48, Iter: 58600, Loss: 106.8, Recon: 94.16
Epoch: 49, Iter: 58800, Loss: 104.4, Recon: 91.79
Epoch: 49, Iter: 59000, Loss: 115.3, Recon: 102.6
Epoch: 49, Iter: 59200, Loss: 112.7, Recon: 99.65
Epoch: 49, Iter: 59400, Loss: 114.0, Recon: 101.0
Epoch: 49, Iter: 59600, Loss: 111.8, Recon: 98.97
Epoch: 49, Iter: 59800, Loss: 108.1, Recon: 95.31
Epoch: 50, Iter: 60000, Loss: 104.8, Recon: 92.23
Epoch: 50, Iter: 60200, Loss: 114.9, Recon: 102.3
Epoch: 50, Iter: 60400, Loss: 112.3, Recon: 99.15
Epoch: 50, Iter: 60600, Loss: 113.6, Recon: 100.8
Epoch: 50, Iter: 60800, Loss: 111.2, Recon: 98.57
Epoch: 50, Iter: 61000, Loss: 106.6, Recon: 93.99
Epoch: 51, Iter: 61200, Loss: 104.9, Recon: 92.35
Epoch: 51, Iter: 61400, Loss: 114.7, Recon: 102.2
Epoch: 51, Iter: 61600, Loss: 112.7, Recon: 99.57
Epoch: 51, Iter: 61800, Loss: 113.7, Recon: 100.9
Epoch: 51, Iter: 62000, Loss: 111.9, Recon: 99.37
Epoch: 51, Iter: 62200, Loss: 107.6, Recon: 94.75
Epoch: 52, Iter: 62400, Loss: 105.9, Recon: 93.52
Epoch: 52, Iter: 62600, Loss: 115.8, Recon: 103.0
Epoch: 52, Iter: 62800, Loss: 112.5, Recon: 99.31
Epoch: 52, Iter: 63000, Loss: 114.3, Recon: 101.4
Epoch: 52, Iter: 63200, Loss: 110.9, Recon: 98.25
Epoch: 52, Iter: 63400, Loss: 106.6, Recon: 93.83
Epoch: 53, Iter: 63600, Loss: 104.1, Recon: 91.65
Epoch: 53, Iter: 63800, Loss: 114.0, Recon: 101.3
Epoch: 53, Iter: 64000, Loss: 114.0, Recon: 100.7
Epoch: 53, Iter: 64200, Loss: 114.5, Recon: 101.5
Epoch: 53, Iter: 64400, Loss: 112.3, Recon: 99.65
Epoch: 53, Iter: 64600, Loss: 105.9, Recon: 93.17
Epoch: 54, Iter: 64800, Loss: 105.9, Recon: 93.28
Epoch: 54, Iter: 65000, Loss: 115.0, Recon: 102.3
Epoch: 54, Iter: 65200, Loss: 112.5, Recon: 99.45
Epoch: 54, Iter: 65400, Loss: 114.0, Recon: 100.8
Epoch: 54, Iter: 65600, Loss: 110.6, Recon: 98.05
Epoch: 54, Iter: 65800, Loss: 106.2, Recon: 93.65
Epoch: 55, Iter: 66000, Loss: 105.7, Recon: 93.13
Epoch: 55, Iter: 66200, Loss: 115.2, Recon: 102.6
Epoch: 55, Iter: 66400, Loss: 112.4, Recon: 99.18
Epoch: 55, Iter: 66600, Loss: 114.0, Recon: 101.1
Epoch: 55, Iter: 66800, Loss: 110.9, Recon: 98.19
Epoch: 55, Iter: 67000, Loss: 106.2, Recon: 93.36
Epoch: 56, Iter: 67200, Loss: 105.0, Recon: 92.55
Epoch: 56, Iter: 67400, Loss: 115.6, Recon: 102.8
Epoch: 56, Iter: 67600, Loss: 114.0, Recon: 100.8
Epoch: 56, Iter: 67800, Loss: 113.0, Recon: 99.97
Epoch: 56, Iter: 68000, Loss: 110.6, Recon: 97.87
Epoch: 56, Iter: 68200, Loss: 105.7, Recon: 92.89
Epoch: 57, Iter: 68400, Loss: 104.0, Recon: 91.34
Epoch: 57, Iter: 68600, Loss: 114.7, Recon: 101.9
Epoch: 57, Iter: 68800, Loss: 112.8, Recon: 99.73
Epoch: 57, Iter: 69000, Loss: 114.5, Recon: 101.5
Epoch: 57, Iter: 69200, Loss: 111.3, Recon: 98.56
Epoch: 57, Iter: 69400, Loss: 106.6, Recon: 93.92

```
Epoch: 58, Iter: 69600, Loss: 104.5, Recon: 91.84
Epoch: 58, Iter: 69800, Loss: 115.2, Recon: 102.4
Epoch: 58, Iter: 70000, Loss: 113.6, Recon: 100.6
Epoch: 58, Iter: 70200, Loss: 113.4, Recon: 100.5
Epoch: 58, Iter: 70400, Loss: 111.3, Recon: 98.63
Epoch: 58, Iter: 70600, Loss: 107.3, Recon: 94.51
Epoch: 59, Iter: 70800, Loss: 106.1, Recon: 93.43
Epoch: 59, Iter: 71000, Loss: 114.0, Recon: 101.2
Epoch: 59, Iter: 71200, Loss: 112.1, Recon: 98.89
Epoch: 59, Iter: 71400, Loss: 112.6, Recon: 99.64
Epoch: 59, Iter: 71600, Loss: 110.8, Recon: 98.12
Epoch: 59, Iter: 71800, Loss: 105.2, Recon: 92.44
Epoch: 60, Iter: 72000, Loss: 106.5, Recon: 94.02
Epoch: 60, Iter: 72200, Loss: 114.0, Recon: 101.3
Epoch: 60, Iter: 72400, Loss: 112.7, Recon: 99.54
Epoch: 60, Iter: 72600, Loss: 113.3, Recon: 100.2
Epoch: 60, Iter: 72800, Loss: 111.9, Recon: 99.24
Epoch: 60, Iter: 73000, Loss: 106.3, Recon: 93.56
Epoch: 61, Iter: 73200, Loss: 105.4, Recon: 92.72
Epoch: 61, Iter: 73400, Loss: 115.3, Recon: 102.6
Epoch: 61, Iter: 73600, Loss: 113.2, Recon: 100.1
Epoch: 61, Iter: 73800, Loss: 113.1, Recon: 99.99
Epoch: 61, Iter: 74000, Loss: 110.8, Recon: 98.24
Epoch: 61, Iter: 74200, Loss: 106.6, Recon: 93.9
Epoch: 62, Iter: 74400, Loss: 104.6, Recon: 91.87
Epoch: 62, Iter: 74600, Loss: 113.7, Recon: 100.8
Epoch: 62, Iter: 74800, Loss: 112.2, Recon: 99.03
Epoch: 62, Iter: 75000, Loss: 113.3, Recon: 100.3
Epoch: 62, Iter: 75200, Loss: 110.9, Recon: 98.09
Epoch: 62, Iter: 75400, Loss: 105.2, Recon: 92.45
Epoch: 63, Iter: 75600, Loss: 106.1, Recon: 93.47
Epoch: 63, Iter: 75800, Loss: 115.4, Recon: 102.7
Epoch: 63, Iter: 76000, Loss: 112.9, Recon: 99.86
Epoch: 63, Iter: 76200, Loss: 112.8, Recon: 99.75
Epoch: 63, Iter: 76400, Loss: 111.5, Recon: 98.76
Epoch: 63, Iter: 76600, Loss: 105.5, Recon: 92.65
Epoch: 64, Iter: 76800, Loss: 104.7, Recon: 92.12
Epoch: 64, Iter: 77000, Loss: 114.2, Recon: 101.4
Epoch: 64, Iter: 77200, Loss: 112.7, Recon: 99.4
Epoch: 64, Iter: 77400, Loss: 112.9, Recon: 99.94
Epoch: 64, Iter: 77600, Loss: 110.7, Recon: 97.86
Epoch: 64, Iter: 77800, Loss: 106.5, Recon: 93.64
Epoch: 65, Iter: 78000, Loss: 104.5, Recon: 91.84
Epoch: 65, Iter: 78200, Loss: 115.3, Recon: 102.6
Epoch: 65, Iter: 78400, Loss: 112.2, Recon: 98.93
Epoch: 65, Iter: 78600, Loss: 113.2, Recon: 100.3
Epoch: 65, Iter: 78800, Loss: 112.6, Recon: 99.96
Epoch: 65, Iter: 79000, Loss: 106.5, Recon: 93.61
Epoch: 66, Iter: 79200, Loss: 104.5, Recon: 91.85
Epoch: 66, Iter: 79400, Loss: 113.5, Recon: 100.8
Epoch: 66, Iter: 79600, Loss: 113.0, Recon: 99.94
Epoch: 66, Iter: 79800, Loss: 113.9, Recon: 100.8
Epoch: 66, Iter: 80000, Loss: 111.1, Recon: 98.37
Epoch: 66, Iter: 80200, Loss: 105.6, Recon: 92.7
Epoch: 67, Iter: 80400, Loss: 104.1, Recon: 91.5
Epoch: 67, Iter: 80600, Loss: 114.6, Recon: 101.8
Epoch: 67, Iter: 80800, Loss: 113.1, Recon: 99.76
Epoch: 67, Iter: 81000, Loss: 112.9, Recon: 99.85
Epoch: 67, Iter: 81200, Loss: 110.7, Recon: 97.86
```

```
Epoch: 67, Iter: 81200, Loss: 110.7, Recon: 97.92
Epoch: 67, Iter: 81400, Loss: 105.1, Recon: 92.31
Epoch: 68, Iter: 81600, Loss: 104.5, Recon: 91.92
Epoch: 68, Iter: 81800, Loss: 113.8, Recon: 101.1
Epoch: 68, Iter: 82000, Loss: 112.6, Recon: 99.38
Epoch: 68, Iter: 82200, Loss: 113.4, Recon: 100.4
Epoch: 68, Iter: 82400, Loss: 111.1, Recon: 98.38
Epoch: 68, Iter: 82600, Loss: 105.1, Recon: 92.34
Epoch: 69, Iter: 82800, Loss: 103.8, Recon: 91.15
Epoch: 69, Iter: 83000, Loss: 113.9, Recon: 101.1
Epoch: 69, Iter: 83200, Loss: 113.2, Recon: 99.94
Epoch: 69, Iter: 83400, Loss: 112.0, Recon: 99.07
Epoch: 69, Iter: 83600, Loss: 111.7, Recon: 99.0
Epoch: 69, Iter: 83800, Loss: 105.3, Recon: 92.37
Epoch: 70, Iter: 84000, Loss: 104.5, Recon: 91.76
Epoch: 70, Iter: 84200, Loss: 114.0, Recon: 101.1
Epoch: 70, Iter: 84400, Loss: 111.0, Recon: 97.89
Epoch: 70, Iter: 84600, Loss: 113.6, Recon: 100.5
Epoch: 70, Iter: 84800, Loss: 112.0, Recon: 99.29
Epoch: 70, Iter: 85000, Loss: 105.7, Recon: 93.12
Epoch: 71, Iter: 85200, Loss: 104.7, Recon: 91.96
Epoch: 71, Iter: 85400, Loss: 114.0, Recon: 101.3
Epoch: 71, Iter: 85600, Loss: 111.2, Recon: 97.84
Epoch: 71, Iter: 85800, Loss: 112.8, Recon: 99.81
Epoch: 71, Iter: 86000, Loss: 111.1, Recon: 98.36
Epoch: 71, Iter: 86200, Loss: 105.7, Recon: 93.06
Epoch: 72, Iter: 86400, Loss: 105.9, Recon: 93.23
Epoch: 72, Iter: 86600, Loss: 113.7, Recon: 101.0
Epoch: 72, Iter: 86800, Loss: 111.5, Recon: 98.3
Epoch: 72, Iter: 87000, Loss: 112.9, Recon: 99.84
Epoch: 72, Iter: 87200, Loss: 110.3, Recon: 97.59
Epoch: 72, Iter: 87400, Loss: 104.6, Recon: 91.58
Epoch: 73, Iter: 87600, Loss: 105.1, Recon: 92.35
Epoch: 73, Iter: 87800, Loss: 116.0, Recon: 103.3
Epoch: 73, Iter: 88000, Loss: 111.9, Recon: 98.75
Epoch: 73, Iter: 88200, Loss: 114.4, Recon: 101.3
Epoch: 73, Iter: 88400, Loss: 111.3, Recon: 98.54
Epoch: 73, Iter: 88600, Loss: 105.4, Recon: 92.46
Epoch: 74, Iter: 88800, Loss: 104.0, Recon: 91.32
Epoch: 74, Iter: 89000, Loss: 115.0, Recon: 102.3
Epoch: 74, Iter: 89200, Loss: 112.5, Recon: 99.19
Epoch: 74, Iter: 89400, Loss: 113.3, Recon: 100.1
Epoch: 74, Iter: 89600, Loss: 111.2, Recon: 98.35
Epoch: 74, Iter: 89800, Loss: 105.2, Recon: 92.38
Epoch: 75, Iter: 90000, Loss: 104.0, Recon: 91.4
Epoch: 75, Iter: 90200, Loss: 115.4, Recon: 102.6
Epoch: 75, Iter: 90400, Loss: 111.7, Recon: 98.44
Epoch: 75, Iter: 90600, Loss: 113.4, Recon: 100.4
Epoch: 75, Iter: 90800, Loss: 110.9, Recon: 98.2
Epoch: 75, Iter: 91000, Loss: 104.6, Recon: 91.75
Epoch: 76, Iter: 91200, Loss: 104.6, Recon: 91.84
Epoch: 76, Iter: 91400, Loss: 113.5, Recon: 100.7
Epoch: 76, Iter: 91600, Loss: 112.4, Recon: 99.15
Epoch: 76, Iter: 91800, Loss: 113.7, Recon: 100.4
Epoch: 76, Iter: 92000, Loss: 110.5, Recon: 97.74
Epoch: 76, Iter: 92200, Loss: 104.5, Recon: 91.75
Epoch: 77, Iter: 92400, Loss: 104.3, Recon: 91.57
Epoch: 77, Iter: 92600, Loss: 115.7, Recon: 102.8
Epoch: 77, Iter: 92800, Loss: 112.3, Recon: 99.00
```

```
Epoch: 77, Iter: 92000, Loss: 112.3, Recon: 99.09
Epoch: 77, Iter: 93000, Loss: 113.4, Recon: 100.5
Epoch: 77, Iter: 93200, Loss: 110.9, Recon: 97.98
Epoch: 77, Iter: 93400, Loss: 104.1, Recon: 91.36
Epoch: 78, Iter: 93600, Loss: 104.1, Recon: 91.52
Epoch: 78, Iter: 93800, Loss: 113.5, Recon: 100.8
Epoch: 78, Iter: 94000, Loss: 111.6, Recon: 98.26
Epoch: 78, Iter: 94200, Loss: 112.5, Recon: 99.47
Epoch: 78, Iter: 94400, Loss: 110.8, Recon: 98.03
Epoch: 78, Iter: 94600, Loss: 105.2, Recon: 92.31
Epoch: 79, Iter: 94800, Loss: 103.1, Recon: 90.47
Epoch: 79, Iter: 95000, Loss: 113.3, Recon: 100.5
Epoch: 79, Iter: 95200, Loss: 111.9, Recon: 98.59
Epoch: 79, Iter: 95400, Loss: 113.1, Recon: 100.1
Epoch: 79, Iter: 95600, Loss: 110.9, Recon: 97.97
Epoch: 79, Iter: 95800, Loss: 105.6, Recon: 92.9
Epoch: 80, Iter: 96000, Loss: 104.2, Recon: 91.51
Epoch: 80, Iter: 96200, Loss: 113.8, Recon: 101.0
Epoch: 80, Iter: 96400, Loss: 111.6, Recon: 98.28
Epoch: 80, Iter: 96600, Loss: 113.3, Recon: 100.3
Epoch: 80, Iter: 96800, Loss: 110.6, Recon: 97.67
Epoch: 80, Iter: 97000, Loss: 105.1, Recon: 92.2
Epoch: 81, Iter: 97200, Loss: 104.5, Recon: 91.59
Epoch: 81, Iter: 97400, Loss: 113.2, Recon: 100.3
Epoch: 81, Iter: 97600, Loss: 112.2, Recon: 98.89
Epoch: 81, Iter: 97800, Loss: 112.7, Recon: 99.57
Epoch: 81, Iter: 98000, Loss: 110.7, Recon: 97.84
Epoch: 81, Iter: 98200, Loss: 105.1, Recon: 92.17
Epoch: 82, Iter: 98400, Loss: 104.1, Recon: 91.5
Epoch: 82, Iter: 98600, Loss: 114.1, Recon: 101.2
Epoch: 82, Iter: 98800, Loss: 111.0, Recon: 97.71
Epoch: 82, Iter: 99000, Loss: 115.4, Recon: 102.4
Epoch: 82, Iter: 99200, Loss: 110.9, Recon: 98.14
Epoch: 82, Iter: 99400, Loss: 104.6, Recon: 91.75
Epoch: 83, Iter: 99600, Loss: 104.3, Recon: 91.59
Epoch: 83, Iter: 99800, Loss: 113.2, Recon: 100.4
Epoch: 83, Iter: 100000, Loss: 111.3, Recon: 98.19
Epoch: 83, Iter: 100200, Loss: 113.9, Recon: 100.8
Epoch: 83, Iter: 100400, Loss: 111.1, Recon: 98.5
Epoch: 83, Iter: 100600, Loss: 104.7, Recon: 92.0
Epoch: 84, Iter: 100800, Loss: 103.9, Recon: 91.02
Epoch: 84, Iter: 101000, Loss: 113.5, Recon: 100.7
Epoch: 84, Iter: 101200, Loss: 110.3, Recon: 96.99
Epoch: 84, Iter: 101400, Loss: 112.8, Recon: 99.75
Epoch: 84, Iter: 101600, Loss: 110.0, Recon: 97.16
Epoch: 84, Iter: 101800, Loss: 104.8, Recon: 92.03
Epoch: 85, Iter: 102000, Loss: 105.0, Recon: 92.24
Epoch: 85, Iter: 102200, Loss: 113.1, Recon: 100.2
Epoch: 85, Iter: 102400, Loss: 110.1, Recon: 96.85
Epoch: 85, Iter: 102600, Loss: 113.6, Recon: 100.5
Epoch: 85, Iter: 102800, Loss: 109.9, Recon: 97.0
Epoch: 85, Iter: 103000, Loss: 104.8, Recon: 91.86
Epoch: 86, Iter: 103200, Loss: 104.2, Recon: 91.55
Epoch: 86, Iter: 103400, Loss: 112.8, Recon: 99.84
Epoch: 86, Iter: 103600, Loss: 112.2, Recon: 98.83
Epoch: 86, Iter: 103800, Loss: 113.1, Recon: 100.1
Epoch: 86, Iter: 104000, Loss: 110.9, Recon: 98.03
Epoch: 86, Iter: 104200, Loss: 104.9, Recon: 92.03
Epoch: 87, Iter: 104400, Loss: 103.6, Recon: 90.7
```

```
Epoch: 87, Iter: 104600, Loss: 114.0, Recon: 101.1  
Epoch: 87, Iter: 104800, Loss: 111.0, Recon: 97.69  
Epoch: 87, Iter: 105000, Loss: 112.8, Recon: 99.73  
Epoch: 87, Iter: 105200, Loss: 110.6, Recon: 97.77  
Epoch: 87, Iter: 105400, Loss: 104.6, Recon: 91.74  
Epoch: 88, Iter: 105600, Loss: 104.9, Recon: 92.08  
Epoch: 88, Iter: 105800, Loss: 115.3, Recon: 102.6  
Epoch: 88, Iter: 106000, Loss: 110.7, Recon: 97.44  
Epoch: 88, Iter: 106200, Loss: 113.6, Recon: 100.5  
Epoch: 88, Iter: 106400, Loss: 110.4, Recon: 97.63  
Epoch: 88, Iter: 106600, Loss: 104.4, Recon: 91.59  
Epoch: 89, Iter: 106800, Loss: 103.8, Recon: 90.97  
Epoch: 89, Iter: 107000, Loss: 113.5, Recon: 100.6  
Epoch: 89, Iter: 107200, Loss: 110.7, Recon: 97.34  
Epoch: 89, Iter: 107400, Loss: 112.6, Recon: 99.44  
Epoch: 89, Iter: 107600, Loss: 109.9, Recon: 97.02  
Epoch: 89, Iter: 107800, Loss: 105.6, Recon: 92.67  
Epoch: 90, Iter: 108000, Loss: 104.1, Recon: 91.32  
Epoch: 90, Iter: 108200, Loss: 113.2, Recon: 100.3  
Epoch: 90, Iter: 108400, Loss: 111.7, Recon: 98.35  
Epoch: 90, Iter: 108600, Loss: 114.0, Recon: 101.0  
Epoch: 90, Iter: 108800, Loss: 110.0, Recon: 97.13  
Epoch: 90, Iter: 109000, Loss: 104.6, Recon: 91.8  
Epoch: 91, Iter: 109200, Loss: 103.4, Recon: 90.57  
Epoch: 91, Iter: 109400, Loss: 113.3, Recon: 100.4  
Epoch: 91, Iter: 109600, Loss: 110.8, Recon: 97.39  
Epoch: 91, Iter: 109800, Loss: 111.9, Recon: 98.79  
Epoch: 91, Iter: 110000, Loss: 110.3, Recon: 97.37  
Epoch: 91, Iter: 110200, Loss: 103.5, Recon: 90.7  
Epoch: 92, Iter: 110400, Loss: 103.7, Recon: 90.85  
Epoch: 92, Iter: 110600, Loss: 113.2, Recon: 100.3  
Epoch: 92, Iter: 110800, Loss: 111.6, Recon: 98.41  
Epoch: 92, Iter: 111000, Loss: 112.2, Recon: 99.13  
Epoch: 92, Iter: 111200, Loss: 110.0, Recon: 97.24  
Epoch: 92, Iter: 111400, Loss: 104.3, Recon: 91.47  
Epoch: 93, Iter: 111600, Loss: 103.6, Recon: 90.79  
Epoch: 93, Iter: 111800, Loss: 113.3, Recon: 100.5  
Epoch: 93, Iter: 112000, Loss: 110.4, Recon: 96.92  
Epoch: 93, Iter: 112200, Loss: 113.8, Recon: 100.7  
Epoch: 93, Iter: 112400, Loss: 109.7, Recon: 96.76  
Epoch: 93, Iter: 112600, Loss: 104.1, Recon: 91.18  
Epoch: 94, Iter: 112800, Loss: 103.9, Recon: 91.13  
Epoch: 94, Iter: 113000, Loss: 112.6, Recon: 99.86  
Epoch: 94, Iter: 113200, Loss: 111.2, Recon: 97.72  
Epoch: 94, Iter: 113400, Loss: 113.0, Recon: 99.9  
Epoch: 94, Iter: 113600, Loss: 109.9, Recon: 96.97  
Epoch: 94, Iter: 113800, Loss: 104.2, Recon: 91.36  
Epoch: 95, Iter: 114000, Loss: 104.0, Recon: 91.4  
Epoch: 95, Iter: 114200, Loss: 112.4, Recon: 99.42  
Epoch: 95, Iter: 114400, Loss: 110.5, Recon: 97.15  
Epoch: 95, Iter: 114600, Loss: 113.0, Recon: 100.1  
Epoch: 95, Iter: 114800, Loss: 112.0, Recon: 99.12  
Epoch: 95, Iter: 115000, Loss: 104.8, Recon: 91.73  
Epoch: 96, Iter: 115200, Loss: 104.8, Recon: 92.02  
Epoch: 96, Iter: 115400, Loss: 112.5, Recon: 99.73  
Epoch: 96, Iter: 115600, Loss: 110.6, Recon: 97.3  
Epoch: 96, Iter: 115800, Loss: 113.8, Recon: 100.7  
Epoch: 96, Iter: 116000, Loss: 110.2, Recon: 97.22
```




```
Epoch: 96, Iter: 116200, Loss: 104.7, Recon: 91.72
Epoch: 97, Iter: 116400, Loss: 103.4, Recon: 90.56
Epoch: 97, Iter: 116600, Loss: 112.7, Recon: 99.76
```

Visualize generated samples after training

```
Epoch: 97, Iter: 117200, Loss: 110.4, Recon: 97.55
```

```
z_gen = tf.random.normal(shape=[num_to_show, Z_DIM])
x_gen = P(z_gen)
imgs_numpy = tf.nn.sigmoid(x_gen).numpy()
show_images(imgs_numpy)
plt.show()
```



▼ PART II. Compute the inception score for your trained VAE m

In this part, we will quantitatively measure how good your VAE model is.

▼ Train a classifier

We first need to train a classifier.

```
batch_size = 128
```

```
num_classes = 10
epochs = 20

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = tf.keras.utils.to_categorical(y_train, num_classes)
y_test = tf.keras.utils.to_categorical(y_test, num_classes)

model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Dense(512, activation='relu', input_shape=(784,)))
model.add(tf.keras.layers.Dropout(0.2))
model.add(tf.keras.layers.Dense(512, activation='relu'))
model.add(tf.keras.layers.Dropout(0.2))
model.add(tf.keras.layers.Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=tf.keras.optimizers.RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```



```
60000 train samples
10000 test samples
Model: "sequential_7"
```

Layer (type)	Output Shape	Param #
dense_27 (Dense)	(None, 512)	401920
dropout_2 (Dropout)	(None, 512)	0
dense_28 (Dense)	(None, 512)	262656
dropout_3 (Dropout)	(None, 512)	0
dense_29 (Dense)	(None, 10)	5130

```
Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0
```

```
Epoch 1/20
469/469 [=====] - 2s 4ms/step - loss: 0.2480 - accuracy: 0.9240 - va
Epoch 2/20
469/469 [=====] - 2s 4ms/step - loss: 0.1039 - accuracy: 0.9687 - va
Epoch 3/20
469/469 [=====] - 2s 4ms/step - loss: 0.0760 - accuracy: 0.9771 - va
Epoch 4/20
469/469 [=====] - 2s 4ms/step - loss: 0.0618 - accuracy: 0.9815 - va
Epoch 5/20
469/469 [=====] - 2s 4ms/step - loss: 0.0503 - accuracy: 0.9852 - va
Epoch 6/20
469/469 [=====] - 2s 4ms/step - loss: 0.0454 - accuracy: 0.9870 - va
Epoch 7/20
469/469 [=====] - 2s 4ms/step - loss: 0.0392 - accuracy: 0.9883 - va
Epoch 8/20
469/469 [=====] - 2s 4ms/step - loss: 0.0349 - accuracy: 0.9899 - va
Epoch 9/20
469/469 [=====] - 2s 4ms/step - loss: 0.0311 - accuracy: 0.9908 - va
Epoch 10/20
469/469 [=====] - 2s 4ms/step - loss: 0.0281 - accuracy: 0.9915 - va
Epoch 11/20
469/469 [=====] - 2s 4ms/step - loss: 0.0255 - accuracy: 0.9924 - va
Epoch 12/20
469/469 [=====] - 2s 4ms/step - loss: 0.0254 - accuracy: 0.9927 - va
Epoch 13/20
469/469 [=====] - 2s 4ms/step - loss: 0.0235 - accuracy: 0.9933 - va
Epoch 14/20
469/469 [=====] - 2s 4ms/step - loss: 0.0224 - accuracy: 0.9939 - va
```


▼ Verify the trained classifier on the generated samples

Generate samples and visually inspect if the predicted labels on the samples match the actual dig

```
Epoch 17/20

z_gen = tf.random.normal(shape=[num_samples, Z_DIM])
x_gen = P(z_gen)
imgs_numpy = tf.nn.sigmoid(x_gen[:num_to_show]).numpy()
show_images(imgs_numpy)
```

```
plt.show()
```



```
np.argmax(model.predict(tf.nn.sigmoid(x_gen[:20])), axis=-1)
```

```
array([5, 1, 1, 0, 3, 7, 6, 9, 9, 1, 4, 5, 7, 9, 9, 7, 3, 8, 6, 5])
```

▼ Implement the inception score

Implement Equation 1 in the reference [3]. Replace expectation in the equation with empirical average exponentiation at the end. You should get Inception score of at least 9.0.

```
kld_obj = tf.keras.losses.KLDivergence()
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

score = None
p_y_t = model.predict(tf.nn.sigmoid(x_gen))
p_y = np.ones((num_samples, 1)) * np.mean(p_y_t, axis = 0)
score = np.exp(kld_obj(p_y_t, p_y))

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
print('Inception score: {:.4}'.format(score))
```

↳ Inception score: 9.148

▼ Plot the histogram of predicted labels

Let's additionally inspect the class diversity of the generated samples.

```
plt.hist(np.argmax(model.predict(tf.nn.sigmoid(x_gen)), axis=-1),  
         bins=np.arange(11)-0.5, rwidth=0.8, density=True)  
plt.xticks(range(10))  
plt.show()
```

